

Recommender Systems

MLSS '14

Collaborative Filtering and other approaches

*Xavier Amatriain
Research/Engineering Director @ Netflix*



Index

- 1. Introduction: What is a Recommender System**
- 2. “Traditional” Methods**
 - 2.1. Collaborative Filtering
 - 2.2. Content-based Recommendations
- 3. Novel Methods**
 - 3.1. Learning to Rank
 - 3.2. Context-aware Recommendations
 - 3.2.1. Tensor Factorization
 - 3.2.2. Factorization Machines
 - 3.3. Deep Learning
 - 3.4. Similarity
 - 3.5. Social Recommendations
- 4. Hybrid Approaches**
- 5. A practical example: Netflix**
- 6. Conclusions**
- 7. References**



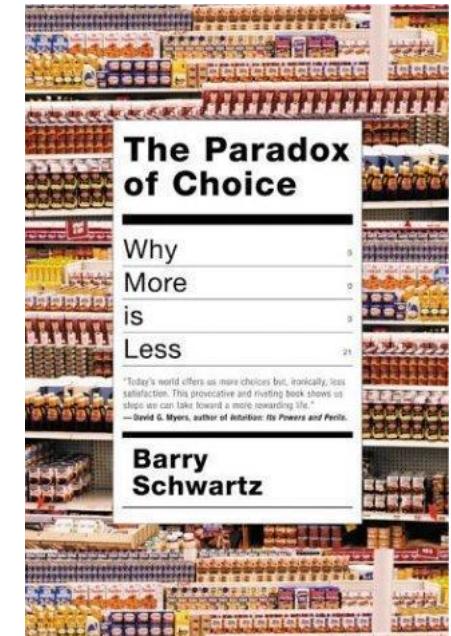
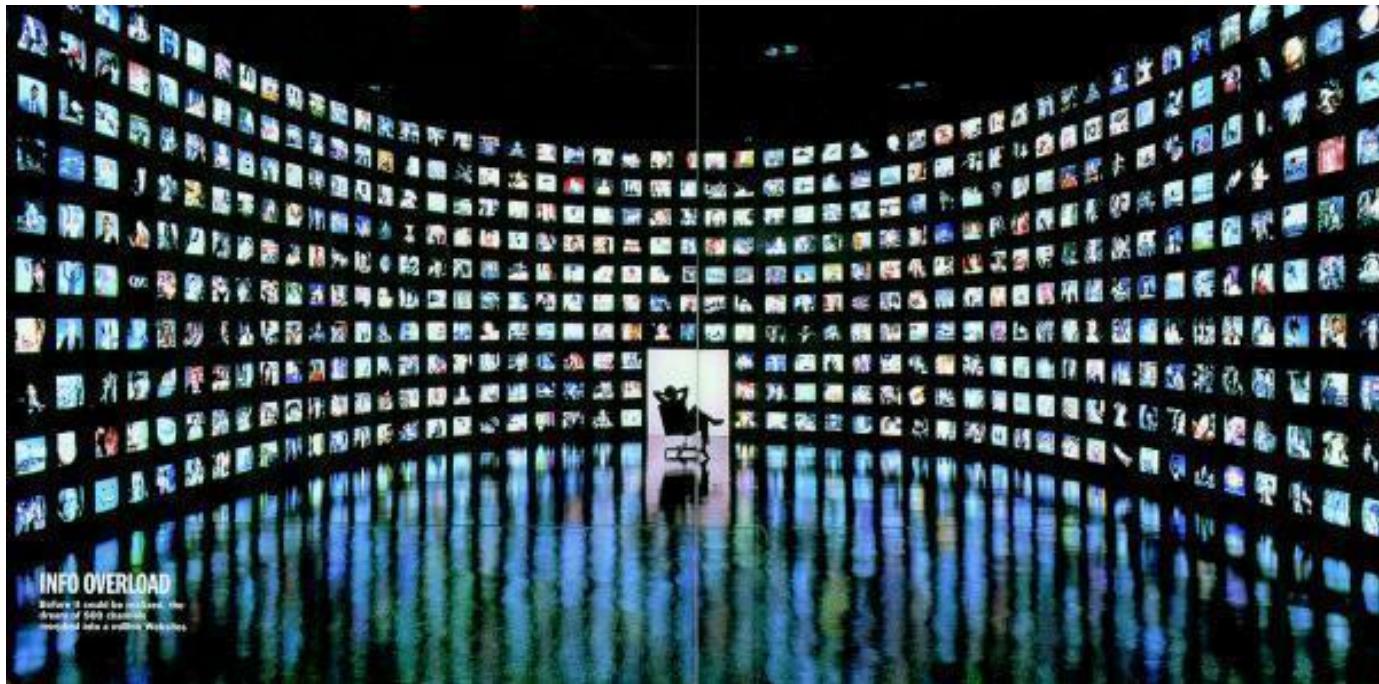
1. Introduction: What is a Recommender System?

The Age of Search has come to an end

- ... long live the **Age of Recommendation!**
- Chris Anderson in “The Long Tail”
 - “*We are leaving the age of information and entering the age of recommendation*”
- CNN Money, “The race to create a 'smart' Google”:
 - “*The Web, they say, is leaving the era of search and entering one of discovery. What's the difference? Search is what you do when you're looking for something. Discovery is when something wonderful that you didn't know existed, or didn't know how to ask for, finds you.*”



Information overload



“People read around 10 MB worth of material a day, hear 400 MB a day, and see 1 MB of information every second” - The Economist, November 2006

In 2015, consumption will raise to 74 GB a day - UCSD Study 2014



The value of recommendations

- Netflix: 2/3 of the movies watched are recommended
- Google News: recommendations generate 38% more clickthrough
- Amazon: 35% sales from recommendations
- Choicestream: 28% of the people would buy more music if they found what they liked.

[Japan Halts US Beef Imports After Banned Meat Found \(Update\)](#)
Bloomberg - 1 hour ago
Jan. 20 (Bloomberg) -- Japan stopped imports of beef from the US after inspectors found banned cattle parts in a shipment, disrupting trade that resumed last month following a two-year halt because of mad-cow disease. ...
[Japan halts US beef imports due to fears of mad cow](#) San Diego Union Tribune
US to probe beef shipment to Japan San Jose Mercury News
[Boston Globe](#) - [Guardian Unlimited](#) - [MarketWatch](#) - [CNN](#) - [all 1,045 related »](#)

[Recommended for gprice@gmail.com »](#) [Learn more](#)

[Serena in denial over her terminal decline](#)
Guardian Unlimited - 7 hours ago - It was in Australia eight years ago that the Williams sisters were seen competing at the same grand slam for the first...
[International Herald Tribune](#) - [TennisReporters.net](#) - [Forbes](#) - [all 319 related »](#)

[2 dozen hurt in Tel Aviv bombing](#)
San Francisco Chronicle - 20 hours ago - Jerusalem -- At least two dozen Israelis were wounded Thursday when a suicide bomber detonated explosives he was...
[Los Angeles Times](#) - [Detroit Free Press](#) - [San Jose Mercury News](#) - [all 836 related »](#)

[US plans to shift diplomats to developing countries](#)
Boston Globe - Jan 19, 2006 - By Farah Stockman, Globe Staff | January 19, 2006. WASHINGTON - Secretary of State Condoleezza Rice announced...
[International Herald Tribune](#) - [Sydney Morning Herald](#) - [Financial Times](#) - [all 70 related »](#)

[Phone Cancer Link Downplayed](#)
Red Herring - [all 170 related »](#)

['American Idol' Gets a Little Mean](#)
Ceres Courier - [all 575 related »](#)

[Deadline to kill US journalist passes with no news](#)
Khaleej Times - [all 2,853 related »](#)

[From here, Oscar race goes inside Hollywood](#)
Reuters - [all 114 related »](#)

[NASA starry-eyed at comet's samples](#)
Houston Chronicle - [all 219 related »](#)

[REGION: Annan urges Iran to resume talks with EU](#)
Daily Times - [all 1,382 related »](#)

[In The News](#)

Osama bin Laden	Midnight Hour
Albert Brooks	Mustang Sally
Tel Aviv	Air Sahara
Jet Airways	Mehmet Ali Agca
Wilson Pickett	Jill Carroll

John Peel: A Life in Music
Michael Heatley

List Price: £6.99
Our Price: £5.59 & eligible for **Free UK delivery** on orders over £15 with Super Saver Delivery. See [details & conditions](#).
You Save: £1.40 (20%)

Availability: usually dispatched within 24 hours.

27 Used & New from £1.60

[See larger photo](#)

Edition: Paperback

[More Product Details](#)

Perfect Partner
Buy John Peel: A Life in Music with [Margrave Of The Marshes](#) today!

  **Total List Price:** £25.98
Buy Together Today: £16.98

[Buy both now](#)

Customers who bought this item also bought:

- [The Little Book of Wanking: The Definitive Guide to Man's Ultimate Relief](#); Paperback ~ Dick Palmer
- [\(Shag Yourself Slim\) The Most Enjoyable Way to Lose Weight](#); Paperback ~ Imah Goer
- [Grumpy Old Men, the Official Handbook](#); Hardcover ~ Stuart Prebble
- [The Little Book of Minge Topiary](#); Paperback ~ Michael O'Mara Books Ltd

The “Recommender problem”

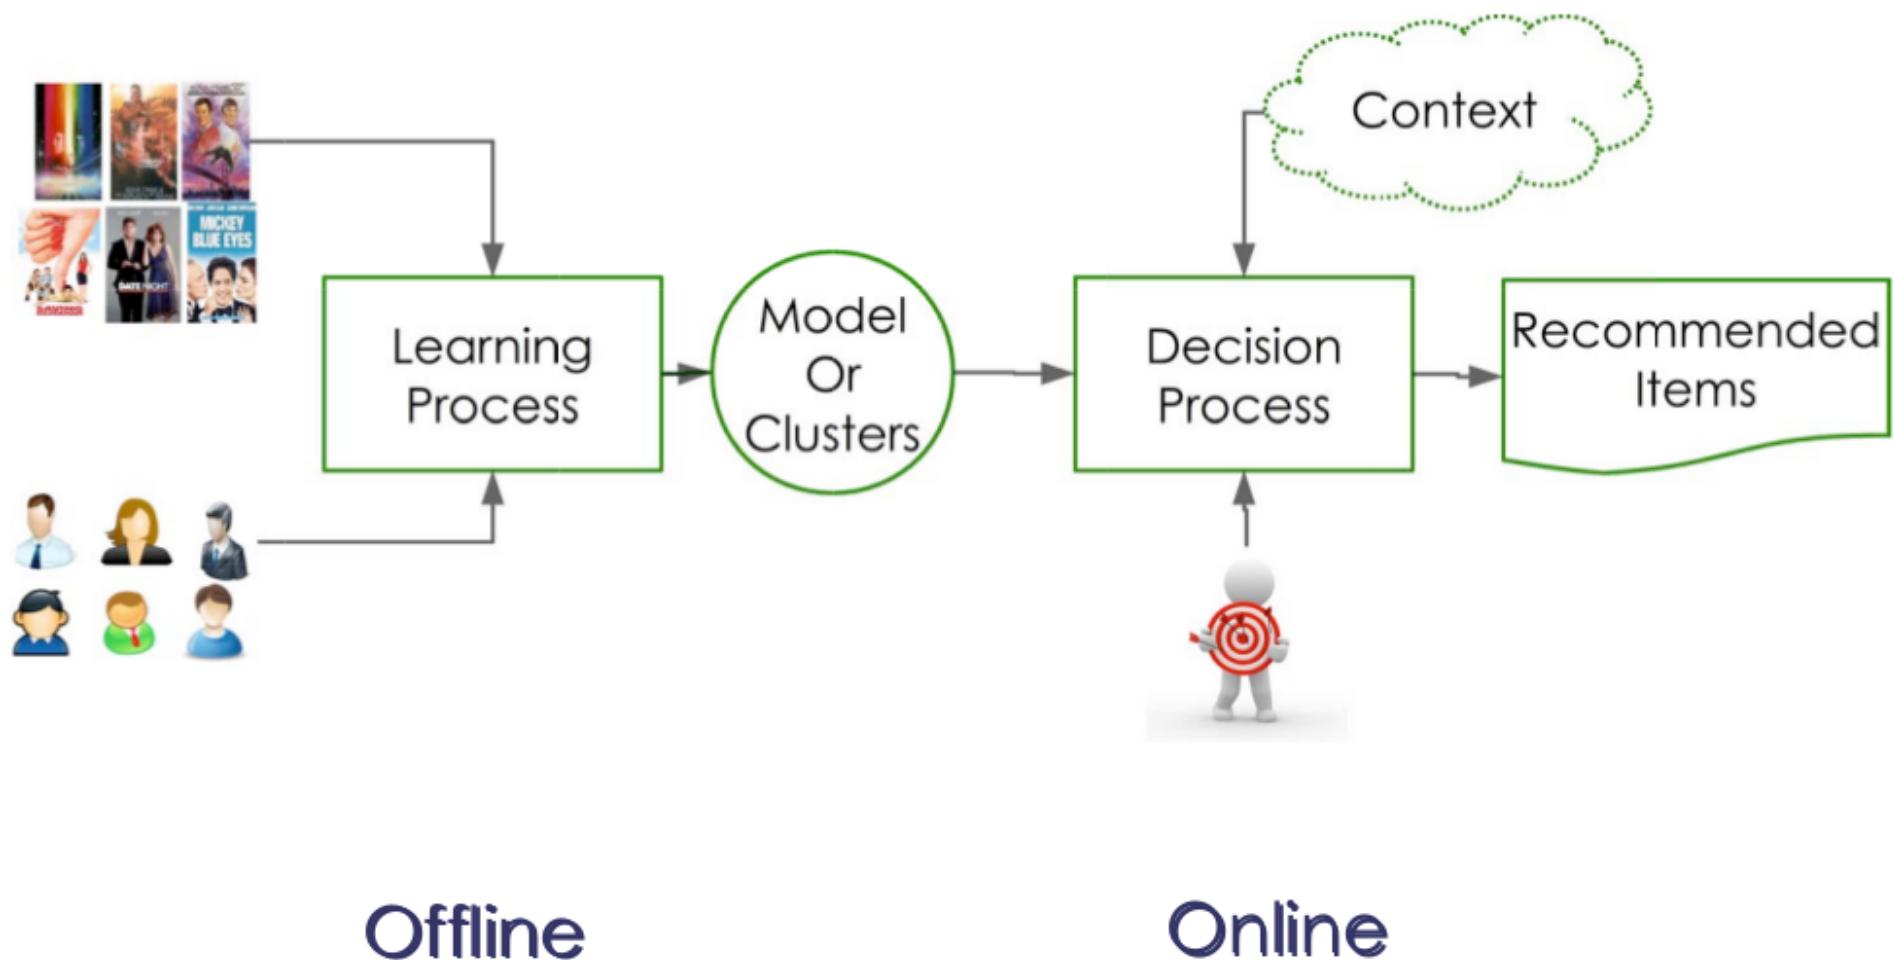
- Estimate a utility function that automatically predicts how a user will like an item.
- Based on:
 - Past behavior
 - Relations to other users
 - Item similarity
 - Context
 - ...

The “Recommender problem”

- Let C be set of all users and let S be set of all possible recommendable items
- Let u be a utility function measuring the usefulness of item s to user c , i.e., $u : C \times S \rightarrow R$, where R is a totally ordered set.
- For each user $c \in C$, we want to choose items $s \in S$ that maximize u .
 - Utility is usually represented by rating but can be any function

$$\forall c \in C, s' = \operatorname{argmax}_c(u(c, s)) \quad s \in S$$

Two-step process



NETFLIX

Approaches to Recommendation

- **Collaborative Filtering:** Recommend items based only on the users past behavior
 - **User-based:** Find similar users to me and recommend what they liked
 - **Item-based:** Find similar items to those that I have previously liked
- **Content-based:** Recommend based on item features
- **Personalized Learning to Rank:** Treat recommendation as a ranking problem
- **Demographic:** Recommend based on user features
- **Social recommendations** (trust-based)
- **Hybrid:** Combine any of the above



What works

- Depends on the **domain** and particular **problem**
- However, in the general case it has been demonstrated that the best isolated approach is CF.
 - Other approaches can be hybridized to improve results in specific cases (cold-start problem...)
- What matters:
 - **Data preprocessing**: outlier removal, denoising, removal of global effects (e.g. individual user's average)
 - “Smart” **dimensionality reduction** using MF/SVD
 - **Combining methods**

Index

1. Introduction: What is a Recommender System
2. “Traditional” Methods
 - 2.1. Collaborative Filtering
 - 2.2. Content-based Recommendations
3. Novel Methods
 - 3.1. Learning to Rank
 - 3.2. Context-aware Recommendations
 - 3.2.1. Tensor Factorization
 - 3.2.2. Factorization Machines
 - 3.3. Deep Learning
 - 3.4. Similarity
 - 3.5. Social Recommendations
4. Hybrid Approaches
5. A practical example: Netflix
6. Conclusions
7. References



2. Traditional Approaches

2.1. Collaborative Filtering

The CF Ingredients

- List of **m Users** and a list of **n Items**
- Each user has a **list of items** with associated **opinion**
 - **Explicit opinion** - a rating score
 - Sometime the rating is **implicitly** – purchase records or listen to tracks
- **Active user** for whom the CF prediction task is performed
- **Metric** for measuring **similarity between users**
- Method for selecting a subset of **neighbors**
- Method for **predicting a rating** for items not currently rated by the active user.

Collaborative Filtering

The basic steps:

1. Identify set of ratings for the **target/active user**
2. Identify set of users most similar to the target/active user according to a similarity function (**neighborhood** formation)
3. Identify the products these similar users liked
4. **Generate a prediction** - rating that would be given by the target user to the product - for each one of these products
5. Based on this predicted rating recommend a set of top N products



Collaborative Filtering

- **Pros:**
 - Requires **minimal knowledge** engineering efforts
 - Users and products are symbols without any internal structure or characteristics
 - Produces good-enough results in most cases
- **Cons:**
 - Requires a large number of **reliable** “user feedback data points” to bootstrap
 - Requires products to be standardized (users should have bought **exactly** the same product)
 - Assumes that **prior behavior determines current behavior** without taking into account “contextual” knowledge (session-level)



Personalised vs Non-Personalised CF

- CF recommendations are **personalized** since the “prediction” is based on the ratings expressed by **similar users**
 - Those **neighbors** are **different** for each target user
- A **non-personalized** collaborative-based recommendation can be generated by averaging the recommendations of **ALL** the users
- How would the two approaches compare?



Personalised vs Non-Personalised CF

Data Set	users	items	total	density	MAE Non Pers	MAE Pers
Jester	48483	100	3519449	0,725	0,220	0,152
MovieLens	6040	3952	1000209	0,041	0,233	0,179
EachMovie	74424	1649	2811718	0,022	0,223	0,151

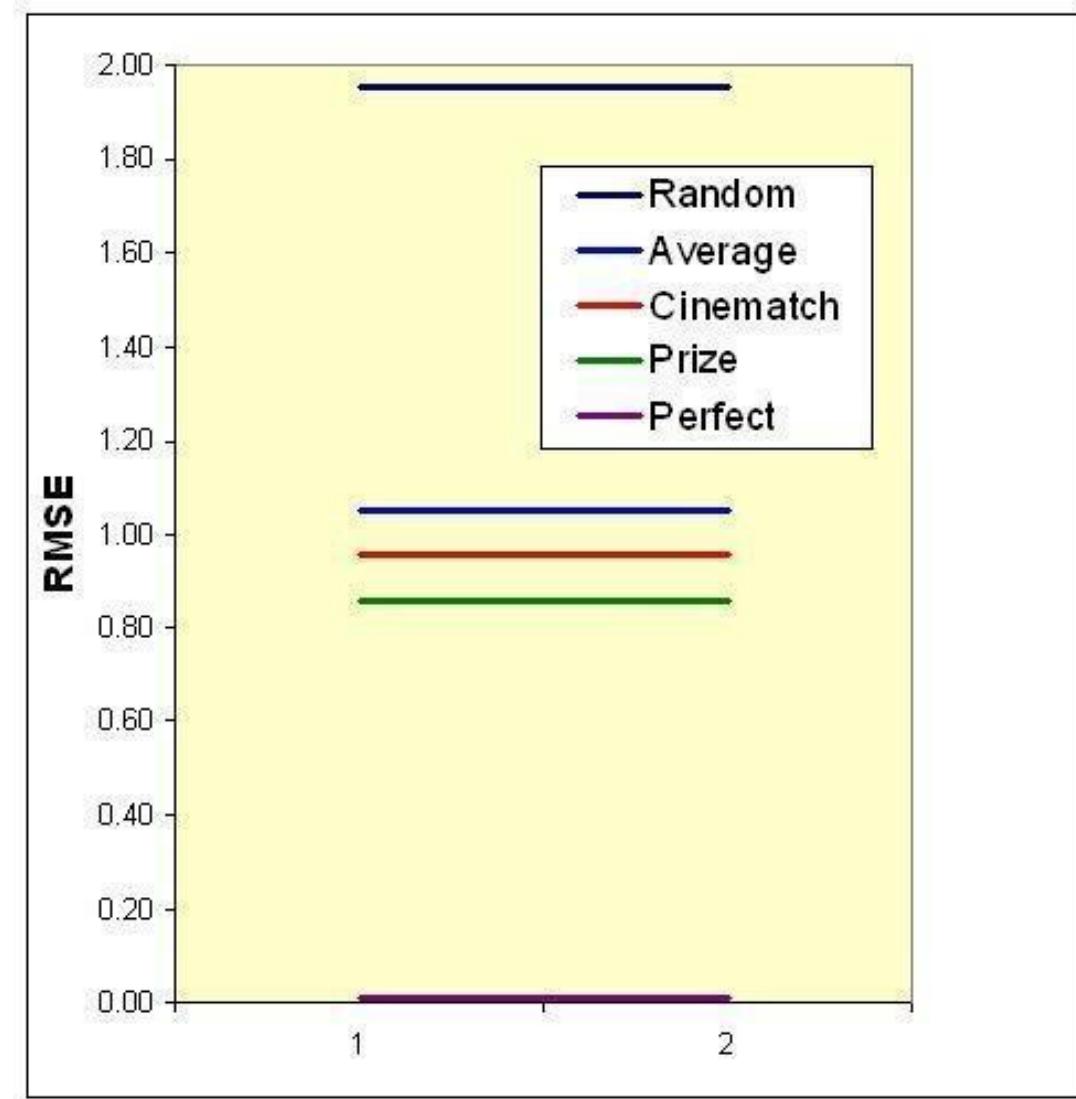
$$MAE_{NP} = \frac{\sum_{i,j} |v_{ij} - v_j|}{num.ratings}$$

v_{ij} is the rating of user i for product j and v_j is the average rating for product j



Personalized vs. Not Personalized

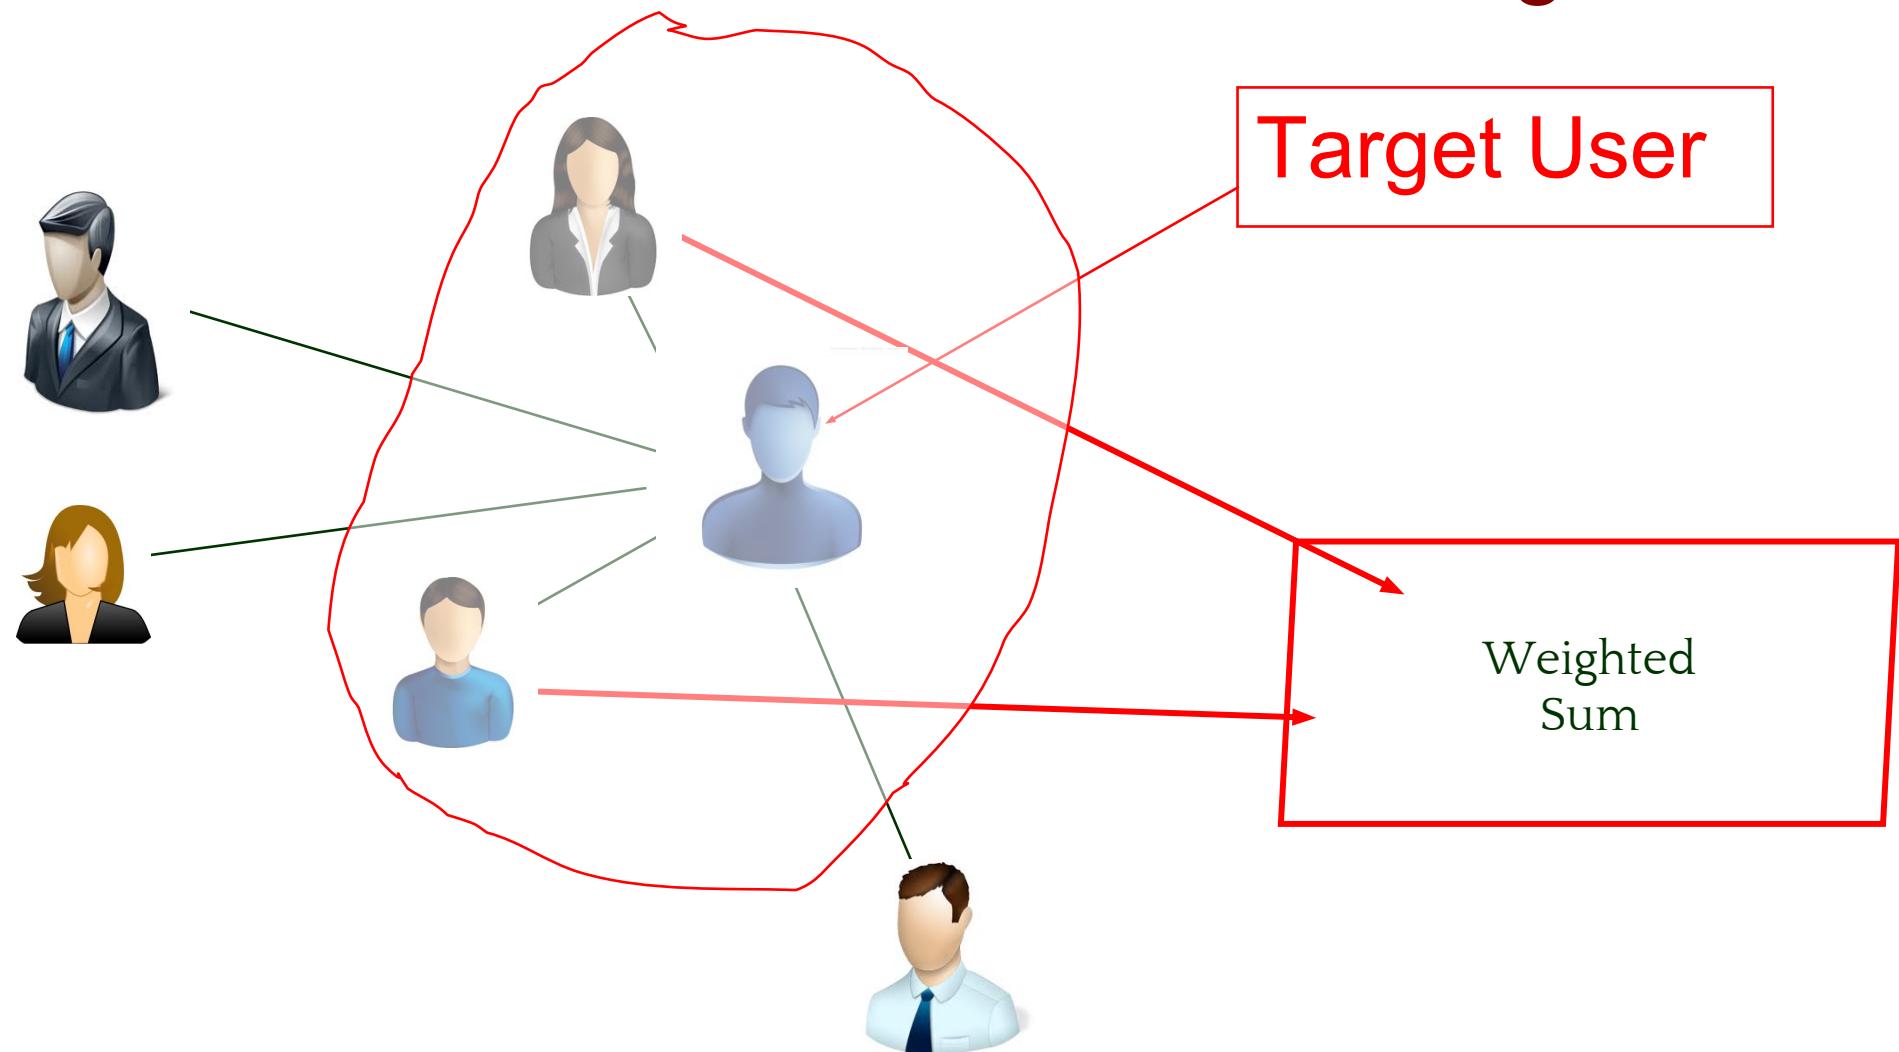
- Netflix Prize's first conclusion: it is really extremely simple to produce “reasonable” recommendations and extremely difficult to improve them.



User-based Collaborative Filtering



User-User Collaborative Filtering



NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

UB Collaborative Filtering

- A collection of user u_i , $i=1, \dots, n$ and a collection of products p_j , $j=1, \dots, m$
- An $n \times m$ matrix of ratings v_{ij} , with $v_{ij} = ?$ if user i did not rate product j
- Prediction for user i and product j is computed

$$v_{ij}^* = K \sum_{v_{kj} \neq ?} u_{jk} v_{kj} \quad \text{or} \quad v_{ij}^* = v_i + K \sum_{v_{kj} \neq ?} u_{jk} (v_{kj} - v_k)$$

- Similarity can be computed by Pearson correlation

$$u_{ik} = \frac{\sum_j (v_{ij} - v_i)(v_{kj} - v_k)}{\sqrt{\sum_j (v_{ij} - v_i)^2 \sum_j (v_{kj} - v_k)^2}} \quad \text{or} \quad \cos(u_i, u_j) = \frac{\sum_{k=1}^m v_{ik} v_{jk}}{\sqrt{\sum_{k=1}^m v_{ik}^2 \sum_{k=1}^m v_{jk}^2}}$$



User-based CF Example



User-based CF Example



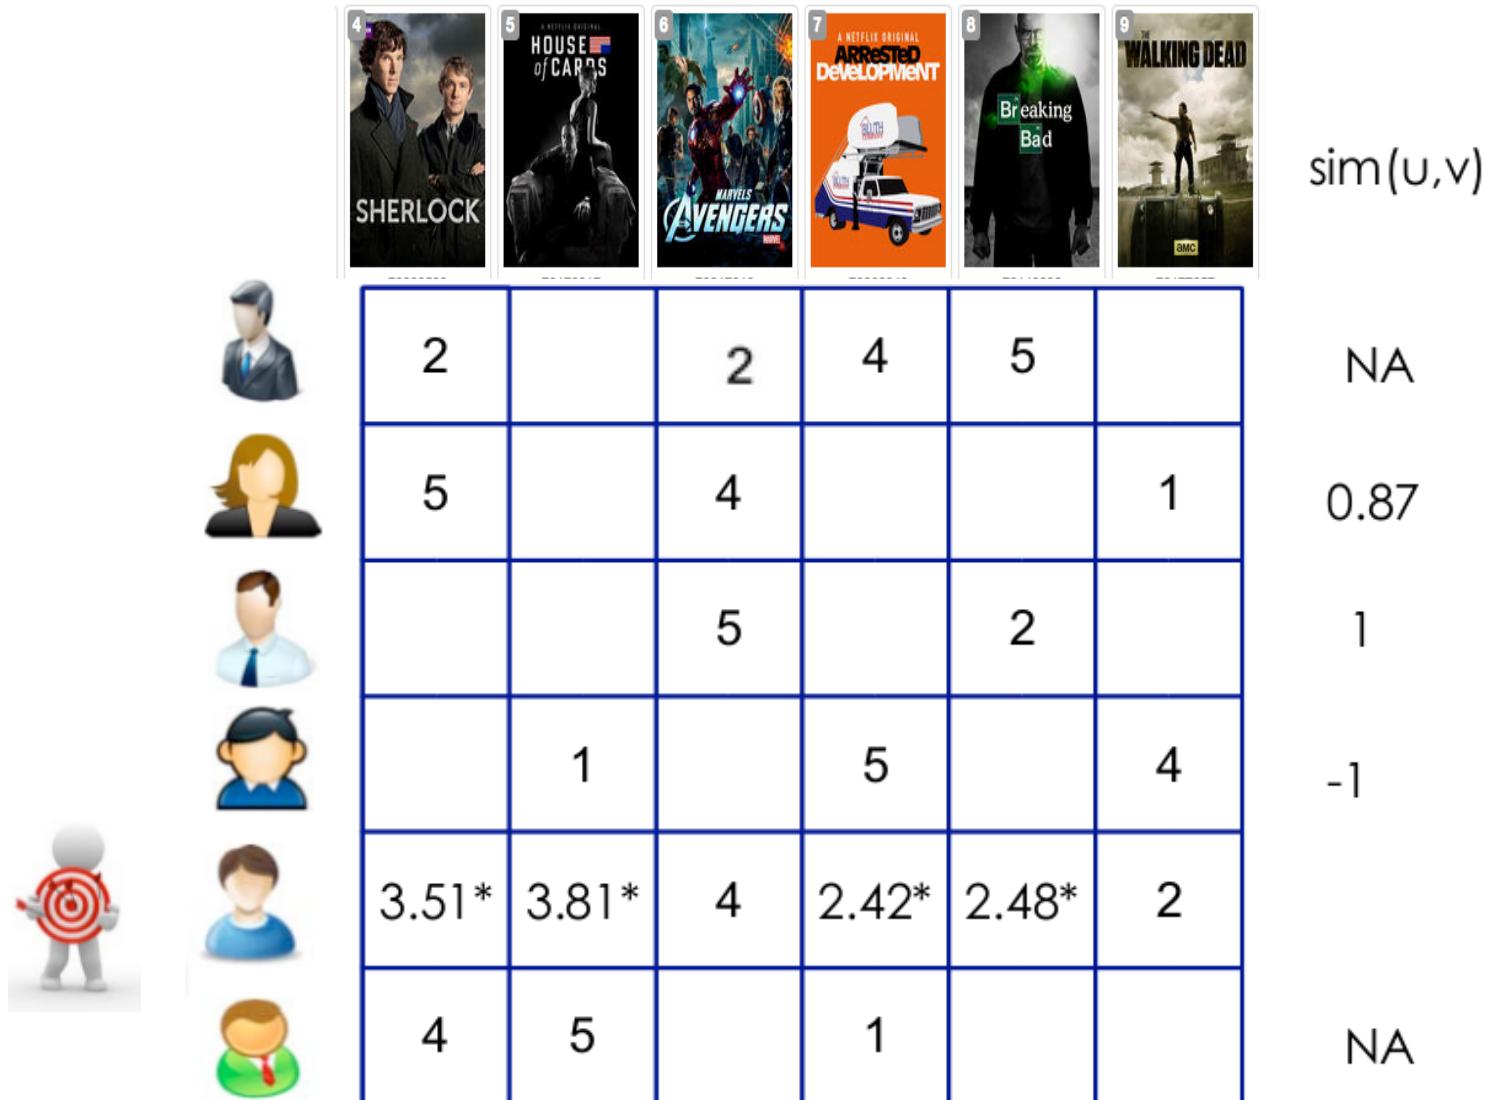
User-based CF Example



User-based CF Example



User-based CF Example



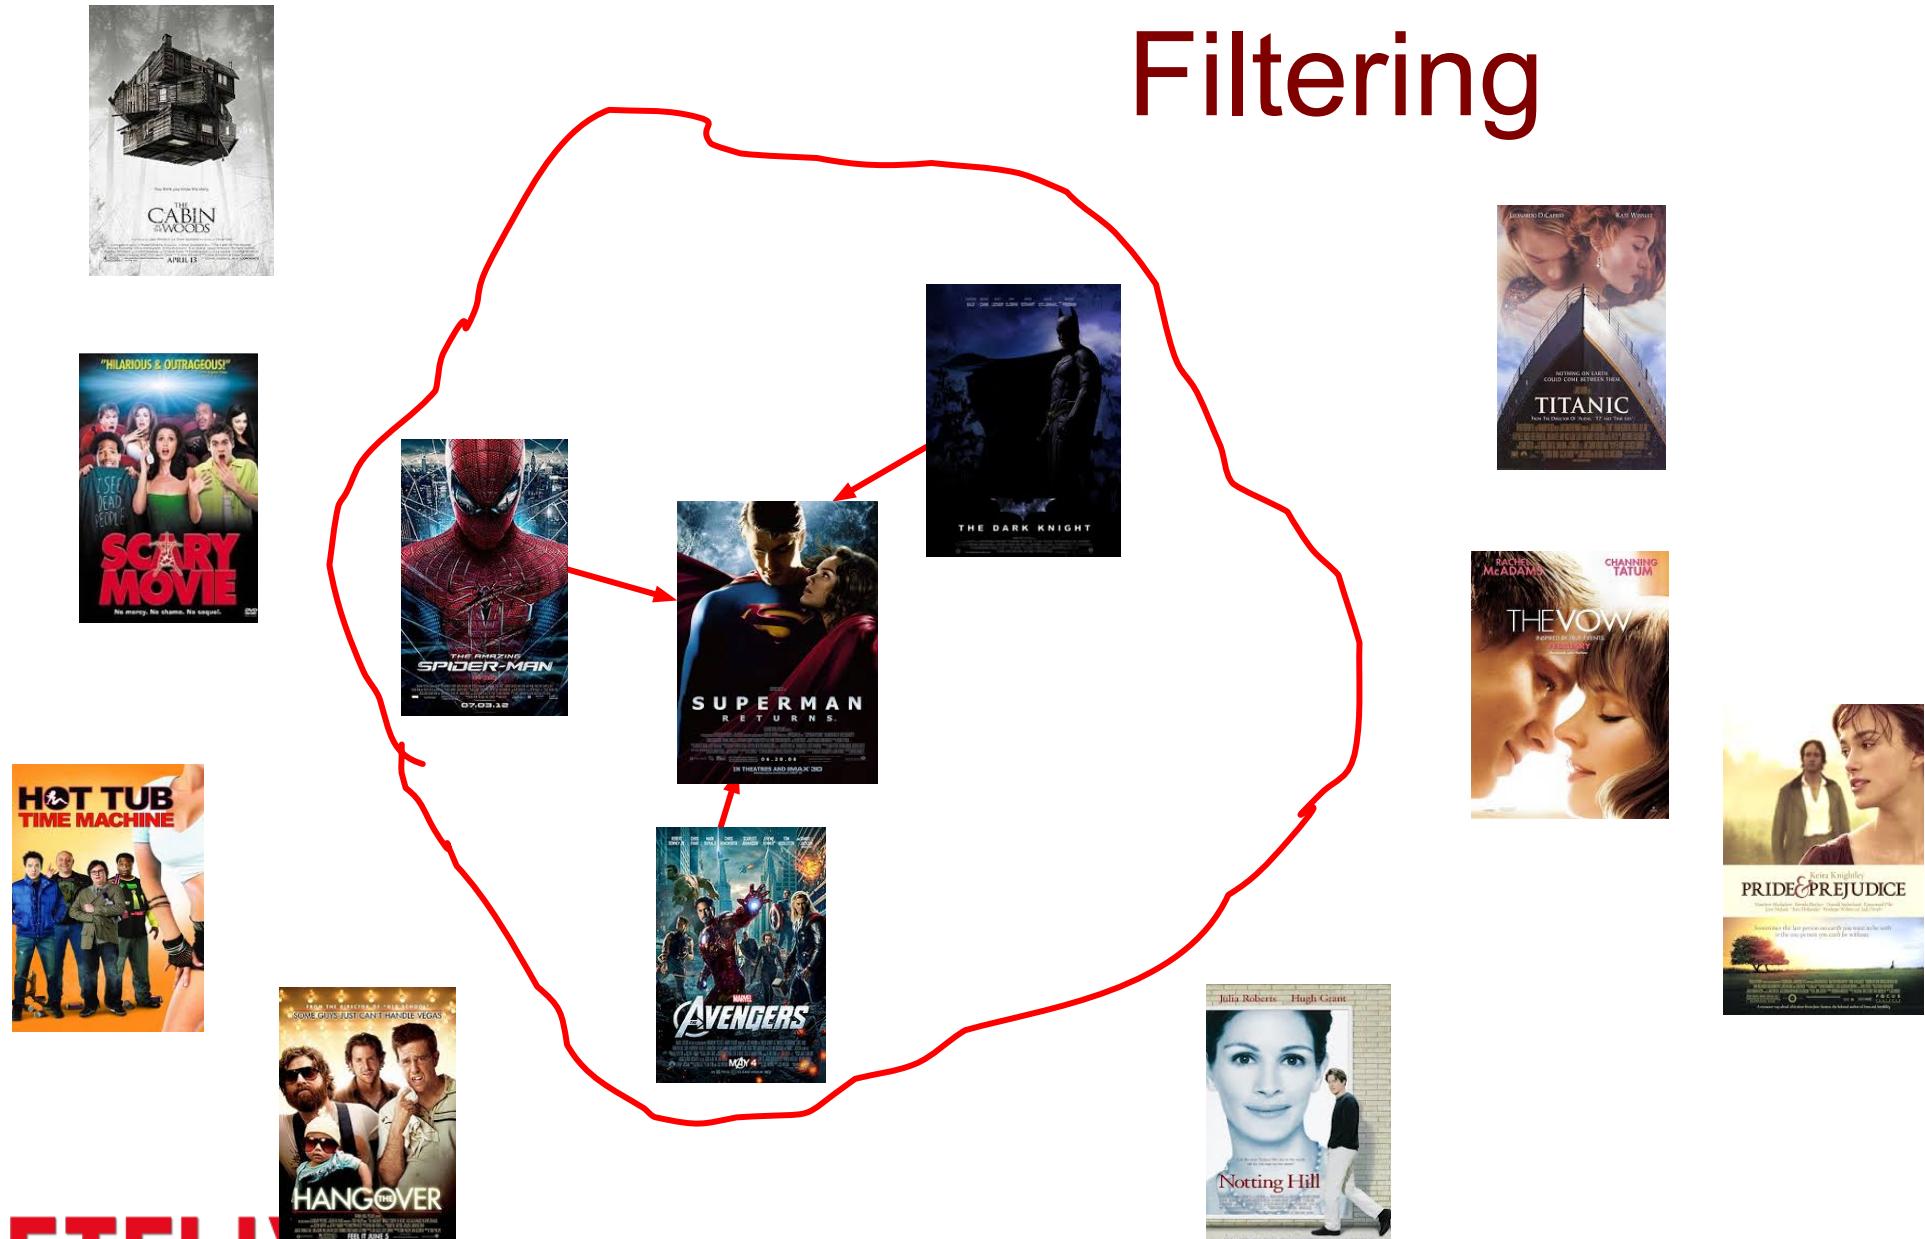
Challenges Of User-based CF Algorithms

- Sparsity – evaluation of large item sets, users purchases are under 1%.
- Difficult to make predictions based on nearest neighbor algorithms => Accuracy of recommendation may be poor.
- Scalability - Nearest neighbor require computation that grows with both the number of users and the number of items.
- Poor relationship among like minded but sparse-rating users.
- Solution : usage of latent models to capture similarity between users & items in a reduced dimensional space.



Item-based Collaborative Filtering

Item-Item Collaborative Filtering



NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

Item Based CF Algorithm

- Look into the items the target user has rated
- Compute how similar they are to the target item
 - Similarity **only using** past **ratings** from other users!
- Select k most similar items.
- Compute Prediction by taking weighted average on the target user's ratings on the most similar items.



Item Similarity Computation

- Similarity between items i & j computed by finding users who have rated them and then applying a similarity function to their ratings.
- Cosine-based Similarity – items are vectors in the m dimensional user space (difference in rating scale between users is not taken into account).

$$S(i, j) = \cos(\vec{i}, \vec{j}) = \frac{\vec{i} \cdot \vec{j}}{\|\vec{i}\|_2 * \|\vec{j}\|_2}$$

Item Similarity Computation

- Correlation-based Similarity - using the Pearson-r correlation (used only in cases where the users rated both item i & item j).

$$S(i, j) = corr_{i,j} = \frac{\sum_{u \in U} (R_{u,i} - \bar{R}_i)(R_{u,j} - \bar{R}_j)}{\sqrt{\sum_{u \in U} (R_{u,i} - \bar{R}_i)^2} \sqrt{\sum_{u \in U} (R_{u,j} - \bar{R}_j)^2}}$$

- $R_{u,i}$ = rating of user u on item i.
- \bar{R}_i = average rating of the i-th item.

Item Similarity Computation

- Adjusted Cosine Similarity – each pair in the co-rated set corresponds to a different user. (takes care of difference in rating scale).

$$S(i, j) = \frac{\sum_{u \in U} (R_{u,i} - \bar{R}_u)(R_{u,j} - \bar{R}_u)}{\sqrt{\sum_{u \in U} (R_{u,i} - \bar{R}_u)^2} \sqrt{\sum_{u \in U} (R_{u,j} - \bar{R}_u)^2}}$$

- $R_{u,i}$ = rating of user u on item i .
- \bar{R}_u = average of the u -th user.



Prediction Computation

- Generating the prediction – look into the target users ratings and use techniques to obtain predictions.
- Weighted Sum – how the active user rates the similar items.

$$P_{u,i} = \frac{\sum_{\text{all similar items}, N} (S_{i,N} * R_{u,N})}{\sum_{\text{all similar items}, N} (|S_{i,N}|)}$$

Item-based CF Example



2			4	5	
5		4			1
	5		2		
	1		5		4
		4			2
4	5		1		

$\text{sim}(i,j)$

-1

NETFLIX

Item-based CF Example



	2		4	5	
	5		4		1
			5		2
		1		5	4
			4		2
	4	5		1	

$\text{sim}(i,j)$

-1 -1



Item-based CF Example



	2			4	5	
	5		4			1
			5		2	
		1		5		4
			4			2
	4	5		1		

$\text{sim}(i,j)$

-1

-1

0.86



Item-based CF Example



2			4	5	
5		4			1
		5		2	
	1		5		4
		4			2
4	5		1		

$\text{sim}(i,j)$

-1

-1

0.86

1



Item-based CF Example



	2			4	5	
	5		4			1
			5		2	
		1		5		4
			4			2
	4	5		1		

sim(i,j) -1 -1 0.86 1 NA

sim(6,5) cannot
be calculated



Item-based CF Example



	2			4	5	2.94*
	5		4			1
			5		2	2.48*
		1		5		4
			4			2
	4	5		1		1.12*

$\text{sim}(i,j)$

-1

-1

0.86

1

NA



Performance Implications

- Bottleneck - Similarity computation.
- Time complexity, highly time consuming with millions of users and items in the database.
 - Isolate the neighborhood generation and predication steps.
 - “off-line component” / “model” – similarity computation, done earlier & stored in memory.
 - “on-line component” – prediction generation process.



Recap: challenges of Nearest-neighbor Collaborative Filtering

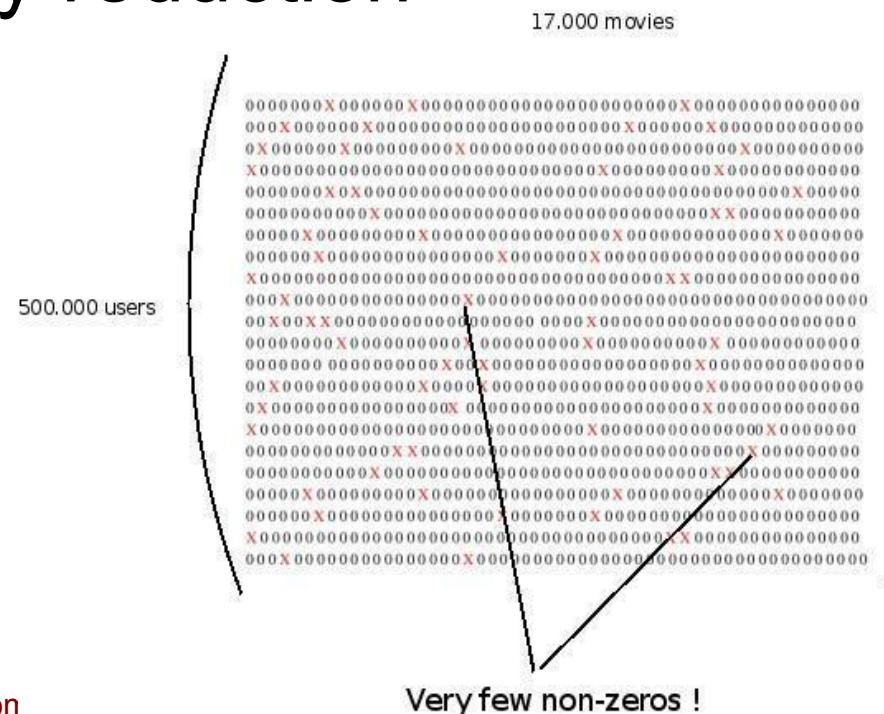


The Sparsity Problem

- Typically: large product sets, user ratings for a small percentage of them
- Example Amazon: millions of books and a user may have bought hundreds of books –
 - the probability that two users that have bought 100 books have a common book (in a catalogue of 1 million books) is 0.01 (with 50 and 10 millions is 0.0002).
- Standard CF must have a number of users comparable to one tenth of the size of the product catalogue

The Sparsity Problem

- If you represent the Netflix Prize rating data in a User/Movie matrix you get...
 - $500,000 \times 17,000 = 8,500 \text{ M positions}$
 - Out of which only 100M are not 0's!
- Methods of dimensionality reduction
 - Matrix Factorization
 - Clustering
 - Projection (PCA ...)



The Scalability Problem

- Nearest neighbor algorithms require computations that grows with both the number of customers and products
- With millions of customers and products a web-based recommender can suffer serious scalability problems
- The worst case complexity is $O(mn)$ (m customers and n products)
- But in practice the complexity is $O(m + n)$ since for each customer only a small number of products are considered
- Some clustering techniques like K-means can help



Performance Implications

- User-based CF – similarity between users is dynamic, precomputing user neighborhood can lead to poor predictions.
- Item-based CF – similarity between items is static.
- enables precomputing of item-item similarity => prediction process involves only a table lookup for the similarity values & computation of the weighted sum.



Other approaches to CF

Model-based Collaborative Filtering



Model Based CF Algorithms

- Memory based
 - Use the entire user-item database to generate a prediction.
 - Usage of statistical techniques to find the neighbors – e.g. nearest-neighbor.
- Model based
 - First develop a model of user
 - Type of model:
 - Probabilistic (e.g. Bayesian Network)
 - Clustering
 - Rule-based approaches (e.g. Association Rules)
 - Classification
 - Regression
 - LDA
 - ...



Model-based CF: What we learned from the Netflix Prize



Netflix Prize

COMPLETED

What we were interested in:

- High quality *recommendations*

Proxy question:

- Accuracy in predicted rating
- Improve by 10% = \$1million!



$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{j=1}^n (y_j - \hat{y}_j)^2}$$

NETFLIX

2007 Progress Prize

- Top 2 algorithms
 - SVD - Prize RMSE: 0.8914
 - RBM - Prize RMSE: 0.8990
- Linear blend Prize RMSE: 0.88
- Currently in use as part of Netflix' rating prediction component
- Limitations
 - Designed for 100M ratings, we have 5B ratings
 - Not adaptable as users add ratings
 - Performance issues



SVD/MF

$$\mathbf{X}[n \times m] = \mathbf{U}[n \times r] \mathbf{S} [r \times r] (\mathbf{V}[m \times r])^\top$$

$$\begin{matrix} X \\ \left(\begin{array}{cccc} x_{11} & x_{12} & \dots & x_{1n} \\ x_{21} & x_{22} & \dots & \\ \vdots & \vdots & \ddots & \\ x_{m1} & & & x_{mn} \end{array} \right) \\ m \times n \end{matrix} = \begin{matrix} U \\ \left(\begin{array}{ccc} u_{11} & \dots & u_{1r} \\ \vdots & \ddots & \\ u_{m1} & & u_{mr} \end{array} \right) \\ m \times r \end{matrix} \begin{matrix} S \\ \left(\begin{array}{ccc} s_{11} & 0 & \dots \\ 0 & \ddots & \\ \vdots & & s_{rr} \end{array} \right) \\ r \times r \end{matrix} \begin{matrix} V^\top \\ \left(\begin{array}{ccc} v_{11} & \dots & v_{1n} \\ \vdots & \ddots & \\ v_{r1} & & v_{rn} \end{array} \right) \\ r \times n \end{matrix}$$

- **X**: $m \times n$ matrix (e.g., m users, n videos)
- **U**: $m \times r$ matrix (m users, r factors)
- **S**: $r \times r$ diagonal matrix (strength of each ‘factor’) (r: rank of the matrix)
- **V**: $r \times n$ matrix (n videos, r factor)

Simon Funk's SVD

- One of the most interesting findings during the Netflix Prize came out of a blog post
- Incremental, iterative, and approximate way to compute the SVD using gradient descent

Monday, December 11, 2006

Netflix Update: Try This at Home



Not to be tied for third place on the [netflix prize](#). And I don't mean a sordid tale of computing in the jungles of the top ten or so.

SVD for Rating Prediction

- User factor vectors $p_u \in \Re^f$ and item-factors vector $q_v \in \Re^f$
- Baseline (bias) $b_{uv} = \mu + b_u + b_v$ (user & item deviation from average)
- Predict rating as $\hat{r}_{uv} = b_{uv} + p_u^T q_v$
- **SVD++** (Koren et. Al) asymmetric variation w. implicit feedback

$$\hat{r}_{uv} = b_{uv} + q_v^T \left(|R(u)|^{-\frac{1}{2}} \sum_{j \in R(u)} (r_{uj} - b_{uj}) x_j + |N(u)|^{-\frac{1}{2}} \sum_{j \in N(u)} y_j \right)$$

- Where
 - $q_v, x_v, y_v \in \Re^f$ are three item factor vectors
 - Users are not parametrized, but rather represented by:
 - $R(u)$: items rated by user u
 - $N(u)$: items for which the user has given implicit preference (e.g. rated vs. not rated)



Clustering

Clustering

- Another way to make recommendations based on past purchases is to **cluster** customers
- Each cluster will be assigned typical preferences, based on preferences of customers who belong to the cluster
- Customers within each cluster will receive recommendations computed at the cluster level

Clustering

	Book1	Book2	Book3	Book4	Book5	Book6
Customer A	X			X		
Customer B		X	X		X	
Customer C		X	X			
Customer D		X				X
Customer E	X				X	

Customers B, C and D are « clustered » together.

Customers A and E are clustered into another separate group

- « Typical » preferences for **CLUSTER** are:
 - Book 2, very high
 - Book 3, high
 - Books 5 and 6, may be recommended
 - Books 1 and 4, not recommended at all

Clustering

	Book1	Book2	Book3	Book4	Book5	Book6
Customer A	X			X		
Customer B		X	X		X	
Customer C		X	X			
Customer D		X				X
Customer E	X				X	
Customer F			X		X	

How does it work?

- Any customer that shall be classified as a member of **CLUSTER** will receive recommendations based on preferences of the group:
 - Book 2 will be highly recommended to Customer F
 - Book 6 will also be recommended to some extent

Clustering

Pros:

- Clustering techniques can be used to work on aggregated data
- Can also be applied as a first step for shrinking the selection of relevant neighbors in a collaborative filtering algorithm and improve performance
- Can be used to capture latent similarities between users or items

Cons:

- Recommendations (per cluster) may be less relevant than collaborative filtering (per individual)



Association Rules

Association rules

- Past purchases are transformed into relationships of common purchases

	Book1	Book2	Book3	Book4	Book5	Book6
Customer A	X			X		
Customer B		X	X		X	
Customer C	.	X	X			.
Customer D		X				X
Customer E	X				X	
Customer F			X		X	

Customers who bought	Also bought ...					
	Book1	Book2	Book3	Book4	Book5	Book6
Book1				1	1	
Book2			2		1	1
Book3		2			2	
Book4	1					
Book5	1	.	1	2		
Book6			1			

Association rules

- These association rules are then used to make recommendations
- If a visitor has some interest in Book 5, she will be recommended to buy Book 3 as well
- Recommendations are constrained to some minimum levels of confidence

		Also bought ...					
		Book1	Book2	Book3	Book4	Book5	Book6
Customers who bought:	Book1				1	1	
	Book2			2		1	1
	Book3		2			2	
	Book4	1					
	Book5	1	1	2			
	Book6		1				



Association rules

Pros:

- Fast to implement (A priori algorithm for frequent itemset mining)
- Fast to execute
- Not much storage space required
- Not « individual » specific
- Very successful in broad applications for large populations, such as shelf layout in retail stores

Cons:

- Not suitable if knowledge of preferences change rapidly
- It is tempting to not apply restrictive confidence rules
→ May lead to literally stupid recommendations



Classifiers

Classifiers

- **Classifiers** are general computational models trained using positive and negative examples
- They may take in inputs:
 - Vector of item features (action / adventure, Bruce Willis)
 - Preferences of customers (like action / adventure)
 - Relations among item
- E.g. Logistic Regression, Bayesian Networks, Support Vector Machines, Decision Trees, etc...

Classifiers

- Classifiers can be used in CF and CB Recommenders
- Pros:
 - Versatile
 - Can be combined with other methods to improve accuracy of recommendations
- Cons:
 - Need a relevant training set
 - May overfit (Regularization)

Limitations of Collaborative Filtering

Limitations of Collaborative Filtering

- **Cold Start:** There needs to be enough other users already in the system to find a match. New items need to get enough ratings.
- **Popularity Bias:** Hard to recommend items to someone with unique tastes.
 - Tends to recommend popular items (items from the tail do not get so much data)



Cold-start

- **New User Problem:** To make accurate recommendations, the system must first learn the user's preferences from the ratings.
 - Several techniques proposed to address this. Most use the hybrid recommendation approach, which combines content-based and collaborative techniques.
- **New Item Problem:** New items are added regularly to recommender systems. Until the new item is rated by a substantial number of users, the recommender system is not able to recommend it.



Index

1. Introduction: What is a Recommender System
2. “Traditional” Methods
 - 2.1. Collaborative Filtering
 - 2.2. Content-based Recommendations
3. Novel Methods
 - 3.1. Learning to Rank
 - 3.2. Context-aware Recommendations
 - 3.2.1. Tensor Factorization
 - 3.2.2. Factorization Machines
 - 3.3. Deep Learning
 - 3.4. Similarity
 - 3.5. Social Recommendations
4. Hybrid Approaches
5. A practical example: Netflix
6. Conclusions
7. References



2.2 Content-based Recommenders



Content-Based Recommendations

- Recommendations based on information on the content of items rather than on other users' opinions/interactions
- Use a machine learning algorithm to induce a model of the users preferences from examples based on a featural description of content.
- In *content-based* recommendations, the system tries to recommend items **similar** to those a given user has liked in the past
- A pure content-based recommender system makes recommendations for a user based solely on the profile built up by **analyzing the content** of items which that user has rated in the past.



What is content?

- What is the **content** of an item?
- It can be explicit **attributes** or **characteristics** of the item. For example for a film:
 - Genre: Action / adventure
 - Feature: Bruce Willis
 - Year: 1995
- It can also be **textual content** (title, description, table of content, etc.)
 - Several techniques to compute the distance between two textual documents
 - Can use NLP techniques to extract content features
- Can be extracted from the signal itself (audio, image)



Content-Based Recommendation

- Common for recommending **text-based products** (web pages, usenet news messages,)
- Items to recommend are “described” by their associated **features** (e.g. keywords)
- **User Model** structured in a “similar” way as the content: features/keywords more likely to occur in the preferred documents (**lazy approach**)
 - Text documents recommended based on a comparison between their content (words appearing) and user model (a set of preferred words)
- The user model can also be a **classifier** based on whatever technique (Neural Networks, Naïve Bayes...)



Advantages of CB Approach

- No need for data on other users.
 - No cold-start or sparsity problems.
- Able to recommend to users with unique tastes.
- Able to recommend new and unpopular items
 - No first-rater problem.
- Can provide explanations of recommended items by listing content-features that caused an item to be recommended.



Disadvantages of CB Approach

- Requires content that can be encoded as meaningful features.
- Some kind of items are not amenable to easy feature extraction methods (e.g. movies, music)
- Even for texts, IR techniques cannot consider multimedia information, aesthetic qualities, download time...
 - If you rate positively a page it could be not related to the presence of certain keywords
- Users' tastes must be represented as a learnable function of these content features.
- Hard to exploit quality judgements of other users.
- Difficult to implement serendipity
- Easy to overfit (e.g. for a user with few data points we may “pigeon hole” her)



Content-based Methods

- Let Content(s) be an *item profile*, i.e. a set of attributes characterizing item s.
- Content usually described with keywords.
- “Importance” (or “informativeness”) of word k_j in document d_j is determined with some weighting measure w_{ij} .
- One of the best-known measures in IR is the term frequency/inverse document frequency (TF-IDF).

$$tf(t, d) = 0.5 + \frac{0.5 \times f(t, d)}{\max\{f(w, d) : w \in d\}} \quad idf(t, D) = \log \frac{N}{|\{d \in D : t \in d\}|}$$

$$tfidf(t, d, D) = tf(t, d) \times idf(t, D)$$



Content-based User Profile

- Let $ContentBasedProfile(c)$ be the profile of user c containing preferences of this user profiles are obtained by:
 - analyzing the content of the previous items
 - using keyword analysis techniques
- For example, $ContentBasedProfile(c)$ can be defined as a vector of weights (w_{c1}, \dots, w_{ck}) , where weight w_{ci} denotes the importance of keyword ki to user c



Similarity Measures

- In content-based systems, the utility function $u(c,s)$ is usually defined as:

$$u(c,s) = \text{score}(\text{ContentBasedProfile}(c), \text{Content}(s))$$

- Both $\text{ContentBasedProfile}(c)$ of user c and $\text{Content}(s)$ of document s can be represented as *TF-IDF* vectors of keyword weights.

Similarity Measurements

- Utility function $u(c,s)$ usually represented by some scoring heuristic defined in terms of vectors , such as the cosine similarity measure.

$$u(c,s) = \cos(\vec{w}_c, \vec{w}_s) = \frac{\vec{w}_c \cdot \vec{w}_s}{\|\vec{w}_c\|_2 \times \|\vec{w}_s\|_2} = \frac{\sum_{i=1}^K w_{i,c} w_{i,s}}{\sqrt{\sum_{i=1}^K w_{i,c}^2} \sqrt{\sum_{i=1}^K w_{i,s}^2}}$$

Statistical and Machine Learning Approaches

Other techniques are feasible

- Bayesian classifiers and various machine learning techniques, including clustering, decision trees, and artificial neural networks.

These methods use models learned from the underlying data rather than heuristics.

- For example, based on a set of Web pages that were rated as “relevant” or “irrelevant” by the user, the naive bayesian classifier can be used to classify unrated Web pages.

Content-based Recommendation. An unrealistic example

- An (unrealistic) example: how to compute recommendations between 8 books based only on their title?
- A customer is interested in the following book: "Building data mining applications for CRM"
- Books selected:
 - Building data mining applications for CRM
 - Accelerating Customer Relationships: Using CRM and Relationship Technologies
 - Mastering Data Mining: The Art and Science of Customer Relationship Management
 - Data Mining Your Website
 - Introduction to marketing
 - Consumer behavior
 - marketing research, a handbook
 - Customer knowledge management



COUNT

	building data mining applications for cm	Accelerating Customer Relationships Using CRM and Relationship Technologies	Mastering Data Mining The Art and Science of Customer Relationship Management	DataMiningYour Website	Introduction to marketing	consumer behavior	marketing research, a handbook	customer knowledge management
a							1	
accelerating		1						
and		1	1					
application	1							
at			1					
behavior						1		
building	1							
consumer						1		
cm	1	1						
customer		1	1					1
data	1		1	1				
for	1							
handbook							1	
introduction					1			
knowledge								1
management			1					1
marketing					1		1	
mastering			1					
mining	1		1	1				
of			1					
relationship		2	1					
research							1	
science			1					
technology		1						
the			1					
to					1			
using		1						
website				1				
your				1				



TFIDF Normed Vectors

	building	customer	data	marketing	research	science	technology	using	website	your	
a	0.502	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
accelerating	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
and	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
application	0.502	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
art	0.000	0.000	0.374	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
behavior	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.707	0.000	0.000	0.000
building	0.502	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
consumer	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.707	0.000	0.000	0.000
crm	0.344	0.296	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
customer	0.000	0.216	0.187	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.381
data	0.251	0.000	0.187	0.316	0.000	0.000	0.000	0.000	0.000	0.000	0.000
for	0.502	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
handbook	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.537	0.000	0.000
introduction	0.000	0.000	0.000	0.000	0.636	0.000	0.000	0.000	0.000	0.000	0.000
knowledge	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.763
management	0.000	0.000	0.256	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.522
marketing	0.000	0.000	0.000	0.000	0.436	0.000	0.000	0.000	0.368	0.000	0.000
mastering	0.000	0.374	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
mining	0.251	0.000	0.187	0.316	0.000	0.000	0.000	0.000	0.000	0.000	0.000
of	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
relationship	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
research	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.537	0.000	0.000
science	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
technology	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
the	0.000	0.000	0.374	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
to	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
using	0.000	0.432	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
website	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000
your	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000	0.000

Mastering Data Mining:
The Art and Science
of Customer Relationship
Management

Data mining
your website

Data

0.187

0.316

Content-based Recommendation

- The system computes distances between this book and the 7 others
- The « closest » books are recommended:
 - **#1:** Data Mining Your Website
 - **#2:** Accelerating Customer Relationships: Using CRM and Relationship Technologies
 - **#3:** Mastering Data Mining: The Art and Science of Customer Relationship Management
 - **Not recommended:** Introduction to marketing
 - **Not recommended:** Consumer behavior
 - **Not recommended:** marketing research, a handbook
 - **Not recommended:** Customer knowledge management



A word of caution

Recommending New Movies: Even a Few Ratings Are More Valuable Than Metadata

István Pilászy *

Dept. of Measurement and Information Systems
Budapest University of Technology and
Economics
Magyar Tudósok krt. 2.
Budapest, Hungary
pila@mit.bme.hu

Domonkos Tikk *,[†]

Dept. of Telecom. and Media Informatics
Budapest University of Technology and
Economics
Magyar Tudósok krt. 2.
Budapest, Hungary
tikk@tmit.bme.hu

ABSTRACT

The Netflix Prize (NP) competition gave much attention to collaborative filtering (CF) approaches. Matrix factorization (MF) based CF approaches assign low dimensional feature vectors to users and items. We link CF and content-based filtering (CBF) by finding a linear transformation that transforms user or item descriptions so that they are as close as possible to the feature vectors generated by MF for CF.

We propose methods for explicit feedback that are able to handle 140 000 features when feature vectors are very sparse. With movie metadata collected for the NP movies we show that the prediction performance of the methods is comparable to that of CF, and can be used to predict user preferences on new movies.

We also investigate the value of movie metadata compared to movie ratings in regards of predictive power. We compare

1. INTRODUCTION

The goal of recommender systems is to give personalized recommendation on items to users. Typically the recommendation is based on the former and current activity of the users, and metadata about users and items, if available.

There are two basic strategies that can be applied when generating recommendations. Collaborative filtering (CF) methods are based only on the activity of users, while content-based filtering (CBF) methods use only metadata. In this paper we propose hybrid methods, which try to benefit from both information sources.

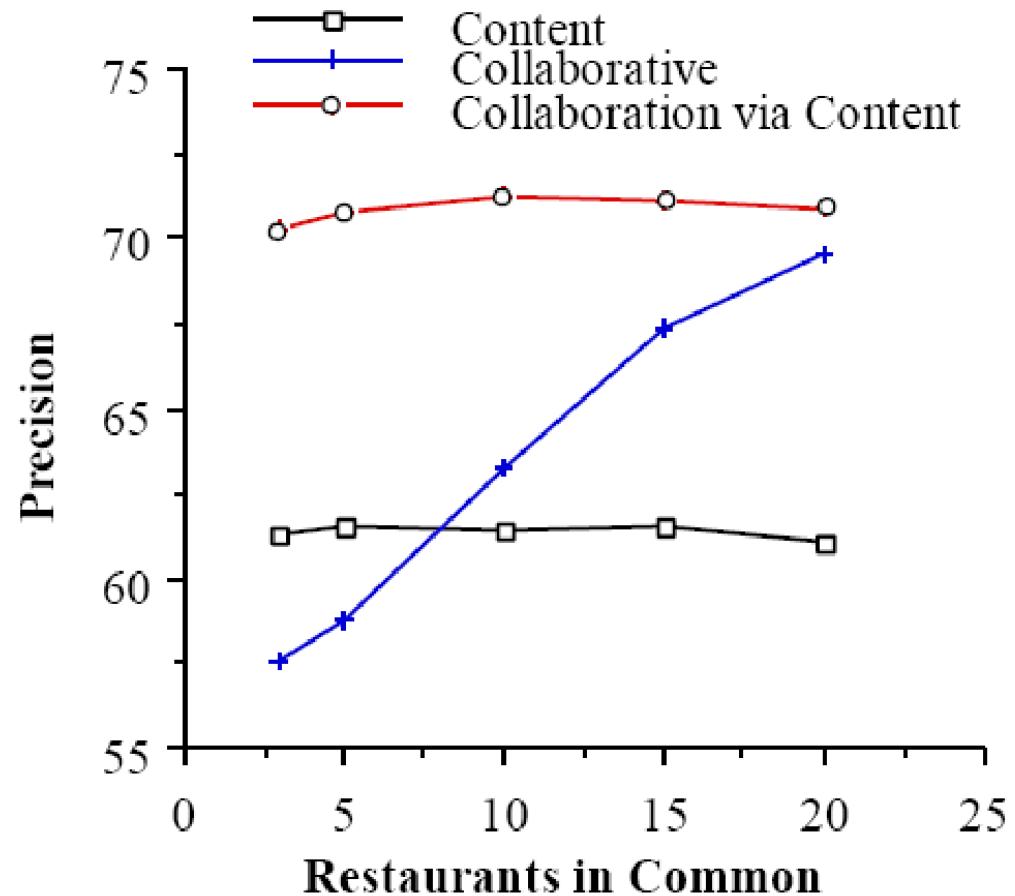
The two most important families of CF methods are matrix factorization (MF) and neighbor-based approaches. Usually, the goal of MF is to find a low dimensional representation for both users and movies, i.e. each user and movie is associated with a feature vector. Movie metadata (which



4 Hybrid Approaches

Comparison of methods (FAB system)

- Content-based recommendation with Bayesian classifier
- Collaborative is standard using Pearson correlation
- Collaboration via content uses the content-based user profiles



Averaged on 44 users

Precision computed in top 3 recommendations

Hybridization Methods

Hybridization Method

Weighted

Description

Outputs from several techniques (in the form of scores or votes) are combined with different degrees of importance to offer final recommendations

Switching

Depending on situation, the system changes from one technique to another

Mixed

Recommendations from several techniques are presented at the same time

Feature combination

Features from different recommendation sources are combined as input to a single technique

Cascade

The output from one technique is used as input of another that refines the result

Feature augmentation

The output from one technique is used as input features to another

Meta-level

The model learned by one recommender is used as input to another



Weighted

- Combine the results of different recommendation techniques into a single recommendation list
 - **Example 1:** a linear combination of recommendation scores
 - **Example 2:** treats the output of each recommender (collaborative, content-based and demographic) as a set of votes, which are then combined in a consensus scheme
- Assumption: relative value of the different techniques is more or less uniform across the space of possible items
 - Not true in general: e.g. a collaborative recommender will be weaker for those items with a small number of raters.



Switching

- The system uses criterion to switch between techniques
 - **Example:** The DailyLearner system uses a content-collaborative hybrid in which a content-based recommendation method is employed first
 - If the content-based system cannot make a recommendation with sufficient confidence, then a collaborative recommendation is attempted
 - Note that switching does not completely avoid the cold-start problem, since both the collaborative and the content-based systems have the “new user” problem
- The main problem of this technique is to identify a GOOD switching condition.



Mixed

- Recommendations from more than one technique are presented together
- The mixed hybrid avoids the “new item” start-up problem
- It does not get around the “new user” start-up problem, since both the content and collaborative methods need some data about user preferences to start up.

Feature Combination

- Features can be combined in several directions. E.g.
 - (1) Treat collaborative information (ratings of users) as additional feature data associated with each example and use content-based techniques over this augmented data set
 - (2) Treat content features as different dimensions for the collaborative setting (i.e. as other ratings from virtual specialized users)

Cascade

- One recommendation technique is employed first to produce a coarse ranking of candidates and a second technique refines the recommendation
 - Example: EntreeC uses its knowledge of restaurants to make recommendations based on the user's stated interests. The recommendations are placed in buckets of equal preference, and the collaborative technique is employed to break ties
- Cascading allows the system to avoid employing the second, lower-priority, technique on items that are already well-differentiated by the first
- But requires a meaningful and constant ordering of the techniques.



Feature Augmentation

- Produce a rating or classification of an item and that information is then incorporated into the processing of the next recommendation technique
 - Example: Libra system makes content-based recommendations of books based on data found in Amazon.com, using a naive Bayes text classifier
 - In the text data used by the system is included “related authors” and “related titles” information that Amazon generates using its internal collaborative systems
- Very similar to the feature combination method:
 - **Here** the output is used for a second RS
 - In **feature combination** the representations used by two systems are combined.



Index

1. Introduction: What is a Recommender System
2. “Traditional” Methods
 - 2.1. Collaborative Filtering
 - 2.2. Content-based Recommendations
3. Novel Methods
 - 3.1. Learning to Rank
 - 3.2. Context-aware Recommendations
 - 3.2.1. Tensor Factorization
 - 3.2.2. Factorization Machines
 - 3.3. Deep Learning
 - 3.4. Similarity
 - 3.5. Social Recommendations
4. Hybrid Approaches
5. A practical example: Netflix
6. Conclusions
7. References



5. Netflix as a practical example



Netflix Prize

COMPLETED

What we were interested in:

- High quality *recommendations*

Proxy question:

- Accuracy in predicted rating
- Improve by 10% = \$1million!

$$\text{RMSE} = \sqrt{\frac{1}{n} \sum_{j=1}^n (y_j - \hat{y}_j)^2}$$

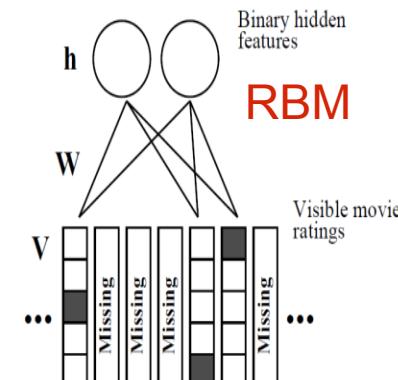
$$\begin{bmatrix} 1 & 1 & 1 & 0 & 0 \\ 2 & 2 & 2 & 0 & 0 \\ 1 & 1 & 1 & 0 & 0 \\ 5 & 5 & 5 & 0 & 0 \\ 0 & 0 & 0 & 2 & 2 \\ 0 & 0 & 0 & 3 & 3 \\ 0 & 0 & 0 & 1 & 1 \end{bmatrix} = \begin{bmatrix} | & | \\ u_1 & u_2 \\ | & | \end{bmatrix} X \begin{bmatrix} \lambda_1 & \emptyset \\ \emptyset & \lambda_2 \end{bmatrix} X \begin{bmatrix} --- & v_1 & --- \\ v_2 & --- \end{bmatrix}$$

SVD



Results

- Top 2 algorithms still in production

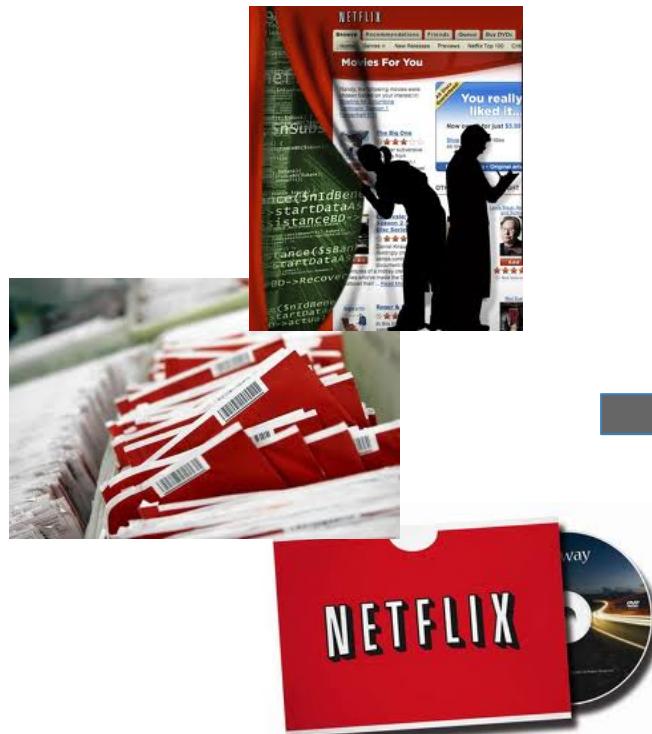


What about the final prize ensembles?

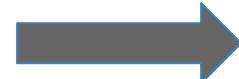
- Our offline studies showed they were too computationally intensive to scale
- Expected improvement not worth the engineering effort
- Plus.... Focus had already shifted to other issues that had more impact than rating prediction.



From the Netflix Prize to today



2006



2014

NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

Anatomy of Netflix Personalization



Everything is a Recommendation





Could Iron Man's Lab Soon Be A Reality?

Photos Photos of this Your Photos Albums Add Contributors

+ Create Album Add Video

Our Trip to Yellowstone

Contributors: Tony Kammann, Kristi Mora and Rader Burns
Contributed about 2 years ago · Taken at Yellowstone National Park (WY)

We decided to go to Yellowstone for the weekend to reconnect with nature. It was a memorable trip with good friends.

Add Photos Tag 6

Facebook To Introduce New Photo Feature

Netflix's New 'My List' Feature Knows You Better Than You Know Yourself (Because Algorithms)

The Huffington Post | By Dino Grandoni

Posted: 08/21/2013 1:44 pm EDT | Updated: 08/22/2013 8:31 am EDT



55 people like this. Be the first of your friends.



Getty

30

12

2

7

107



Share



Tweet



+1



Email



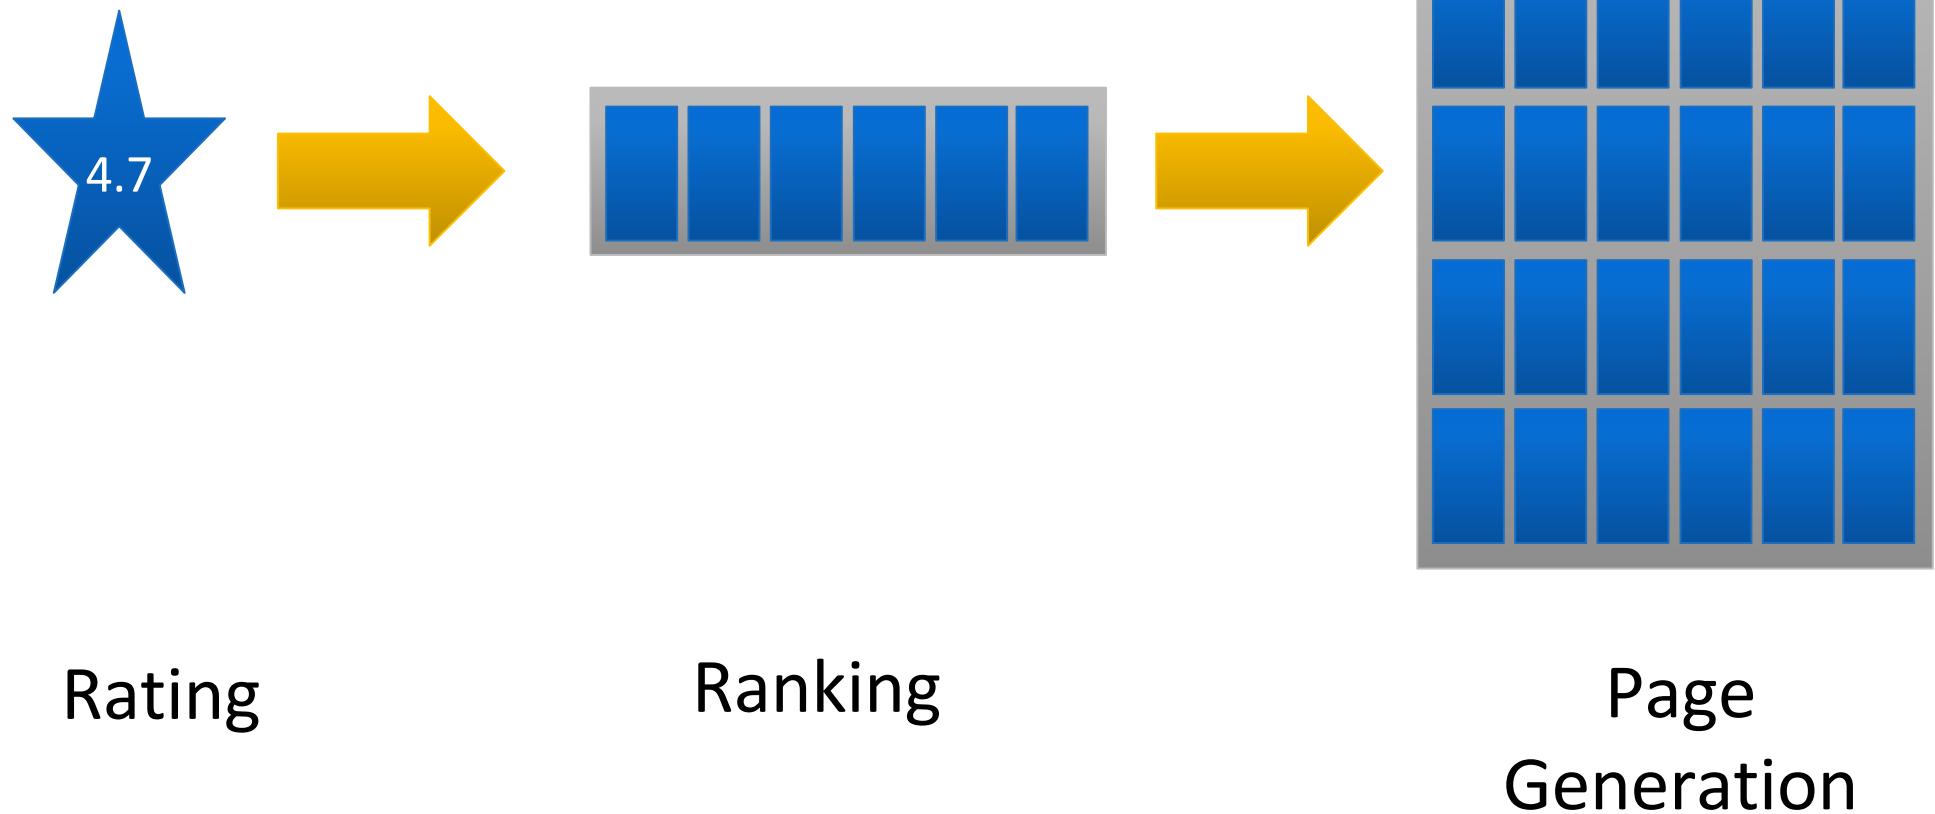
Comment

GET TECHNOLOGY NEWSLETTERS:

Enter email

SUBSCRIBE

Evolution of Recommendation Approach

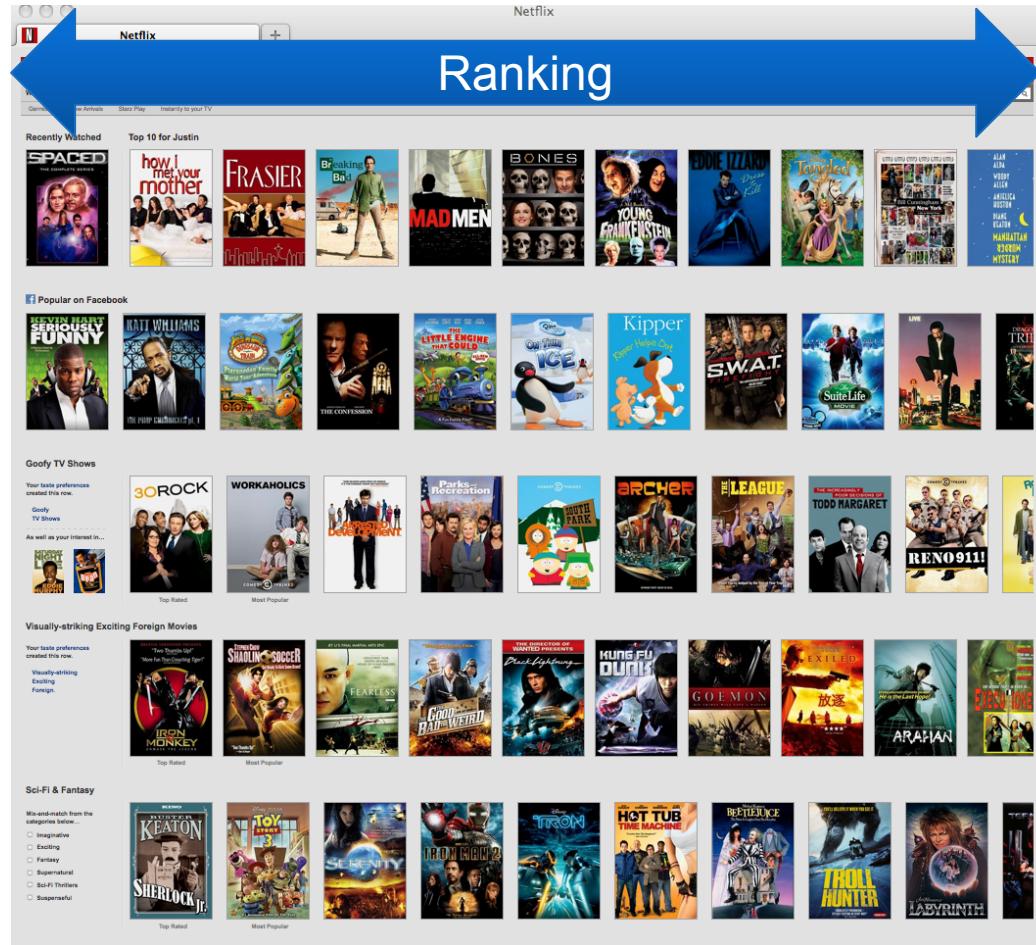


Everything is personalized



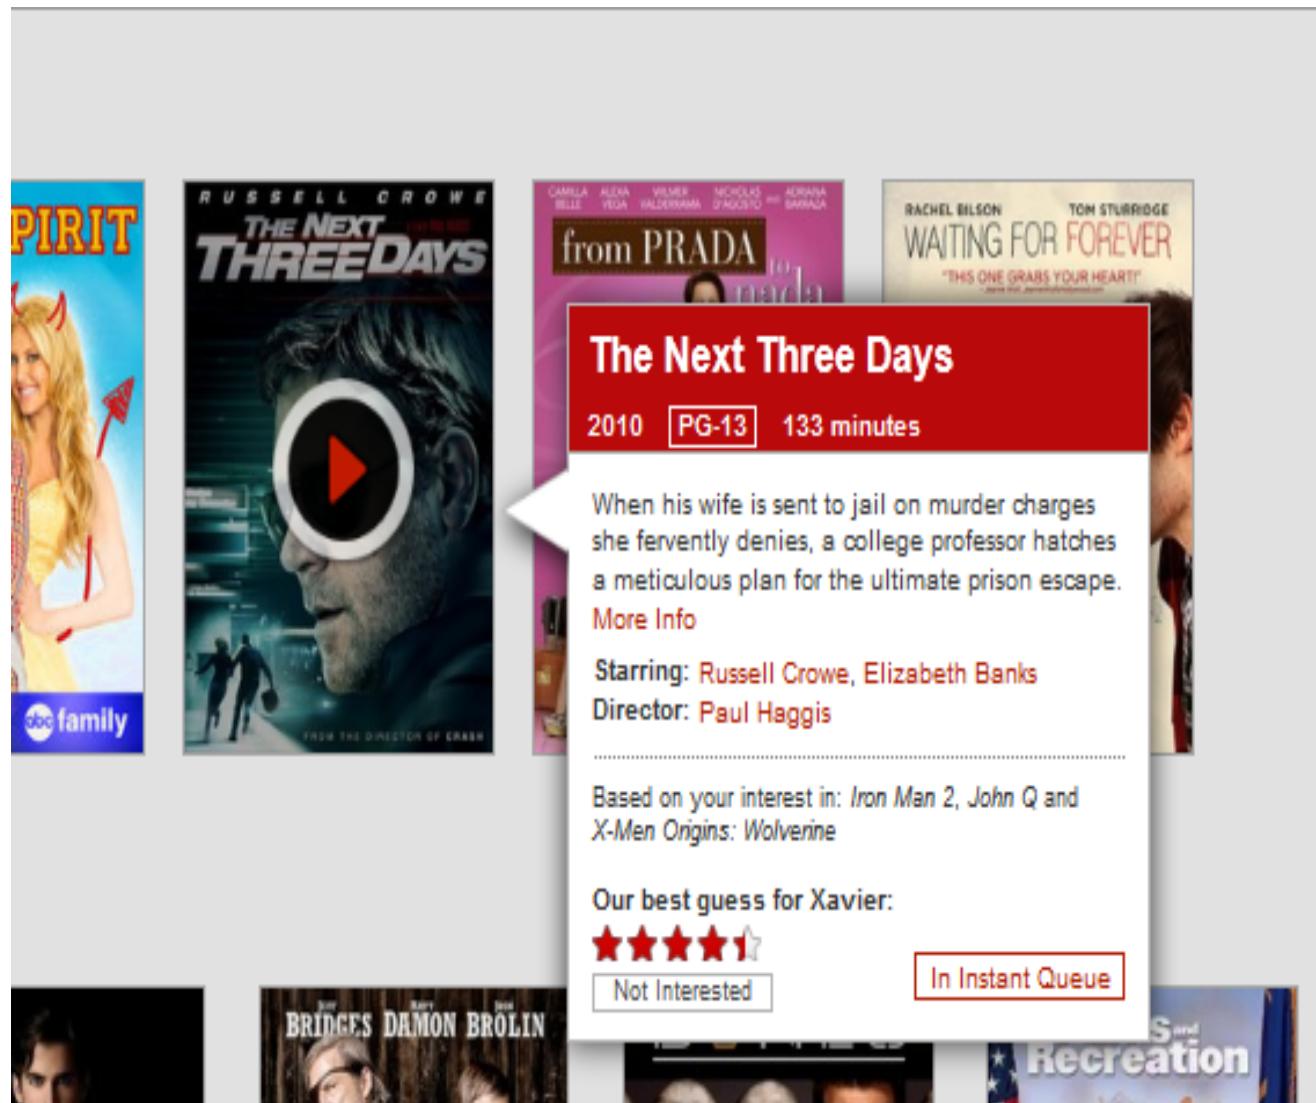
Ranking

Key algorithm, sorts titles in most contexts

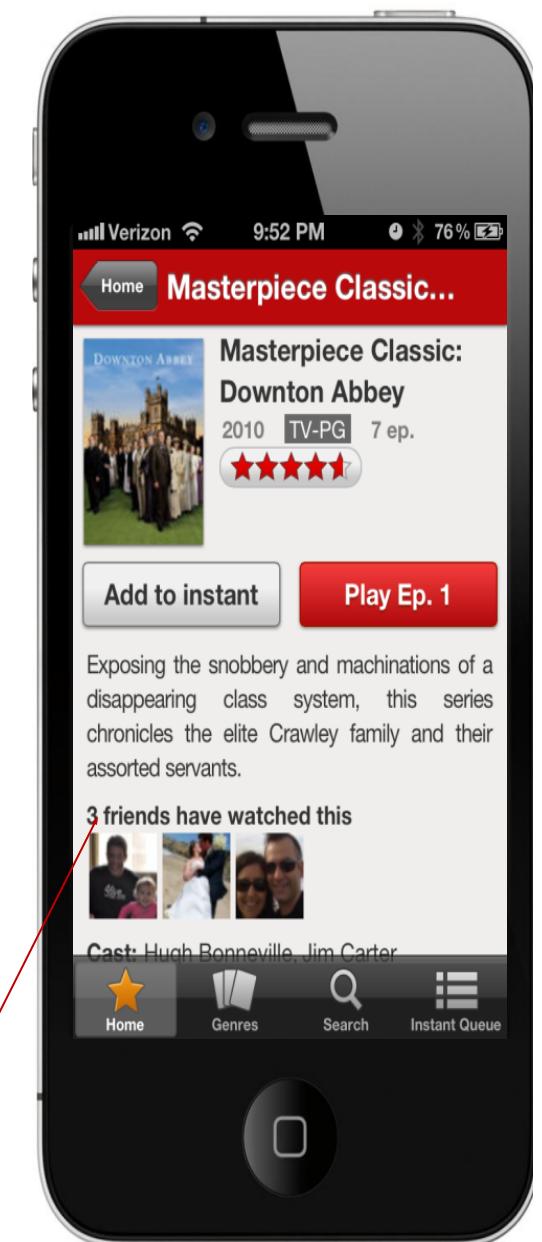


NETFLIX

Support for Recommendations



Social Support

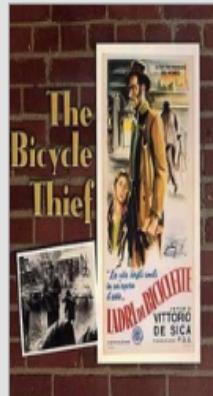
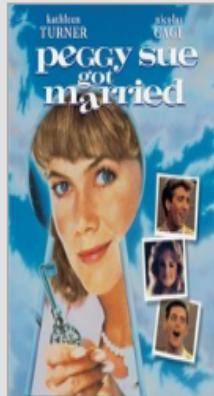
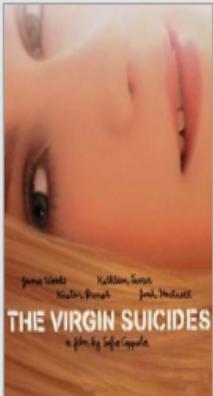


NETFLIX

Social Recommendations

Friends' Favorites

Based on these friends:



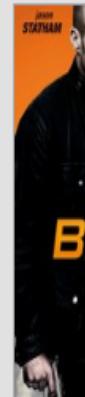
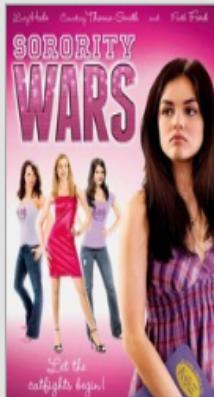
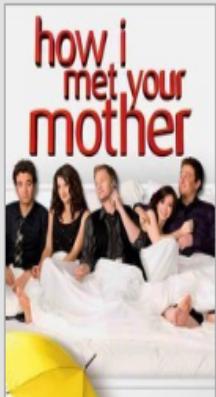
Watched by your friends

Daniel Jacobson

John Ciancutti

Mark White

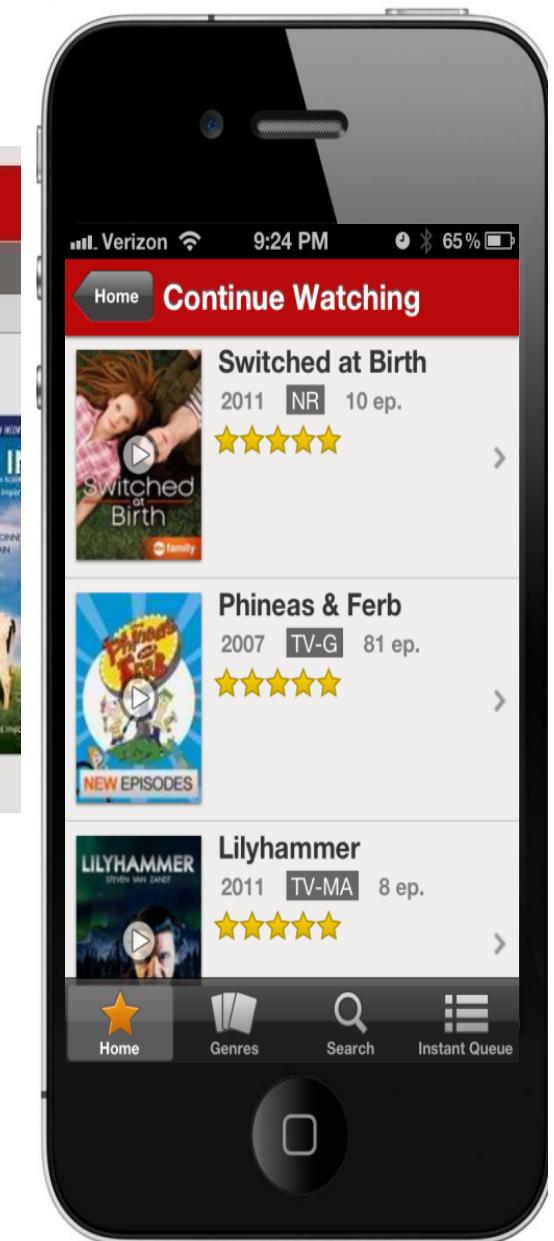
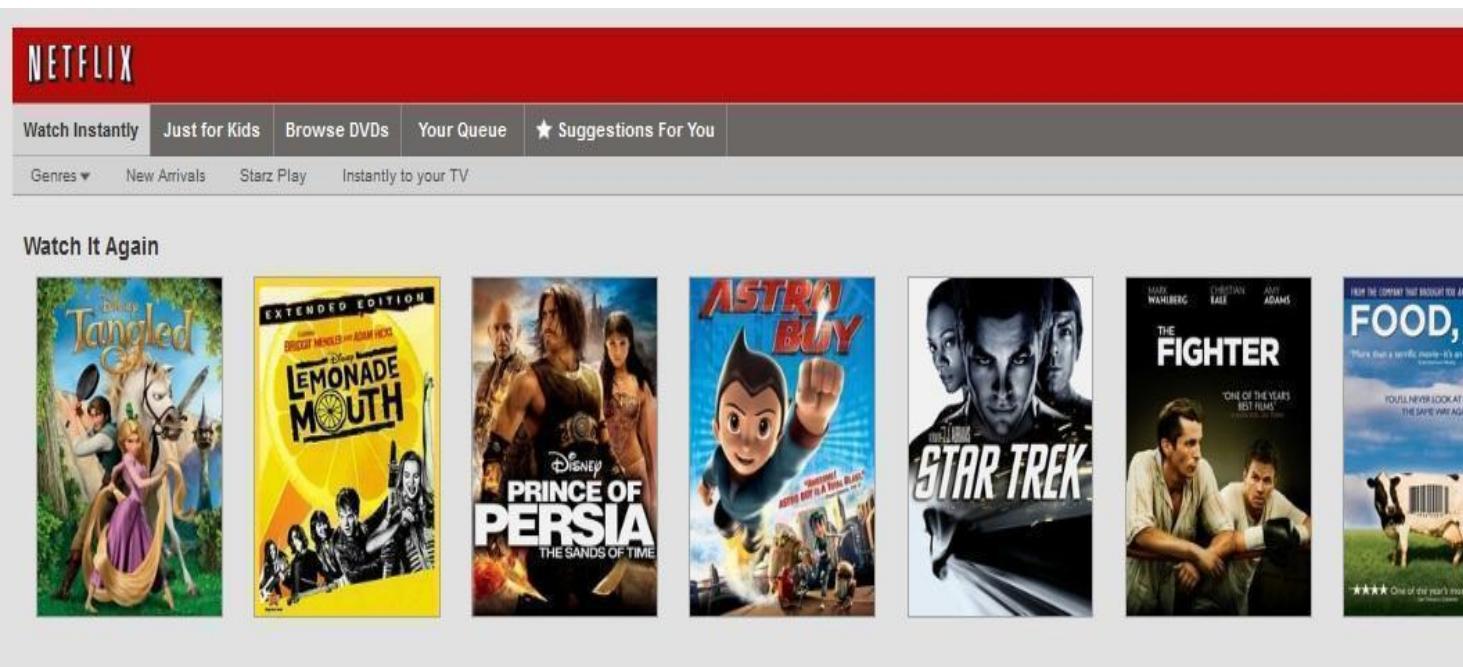
mike Kail



NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

Watch again & Continue Watching



NETFLIX

Genres

Browse

NETFLIX

Independent Underdog Movies

Whale Rider
Good Will Hunting
Billy Elliot
Winning Season
Indian
Bottle Shock
Legendary
Ping Pong Playa

Revenge Movies

Shaolin Soccer
The Punisher
The Girl Who Played with Fire
Accidental Husband
Dragon Tattoo Trilogy
The Red Violin
Blitz
Campus Confidential

Critically-acclaimed Violent Crime Movies

la femme Nikita
Tucker & Dale vs. Evil
Reservoir Dogs
Scarface
RoboCop
Midnight Express
50 Dead Men Walking
Trick 'r Treat

Gory Foreign Movies

Unknown
Resident Evil: Degeneration
Sibyl
9th Company
Dead Snow

NETFLIX

Genre rows

- Personalized genre rows focus on user interest
 - Also provide **context** and “**evidence**”
 - Important for member satisfaction – moving personalized rows to top on devices increased retention
- How are they generated?
 - **Implicit**: based on user's recent plays, ratings, & other interactions
 - **Explicit** taste preferences
 - **Hybrid**: combine the above
- Also take into account:
 - **Freshness** - has this been shown before?
 - **Diversity** – avoid repeating tags and genres, limit number of TV genres, etc.



Genres - personalization

NETFLIX

Watch Instantly Just for Kids Browse DVDs Your Queue ★ Suggestions For You

Suspenseful Wilderness-survival Action & Adventure

Based on your interest in...

Top Rated Most Popular

Independent Dramas Featuring a Strong Female Lead

Your taste preferences created this row.

Independent

As well as your interest in...

Top Rated Most Popular

TV Shows

Mix-and-match from the categories below...

Family-friendly
 TV Comedies
 Cartoons
 Kids' TV Shows
 TV Docs

NET

Genres - personalization

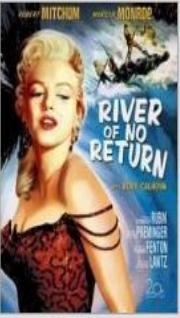
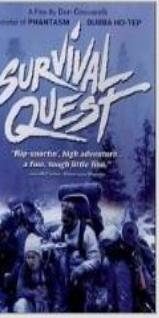
NETFLIX

Watch Instantly Just for Kids Browse DVDs Your Queue ★ Suggestions For You

Genres ▾ New Arrivals Starz Play Instantly to your TV

Suspenseful Wilderness-survival Action & Adventure

Based on your interest in... 

 WILLIAM HOLDEN ROCKY SCHERZINGER
 DENNIS HOPPER
 DENNIS HOPPER
 DENNIS HOPPER
 DENNIS HOPPER
 DENNIS HOPPER
 DENNIS HOPPER

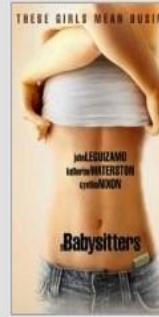
Top Rated Most Popular

Independent Dramas Featuring a Strong Female Lead

Your taste preferences created this row.

Independent

As well as your interest in... 

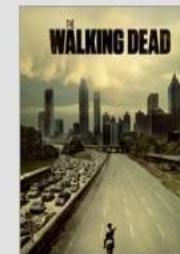
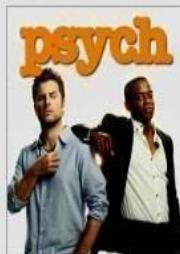
 PRECIOUS
 WINTER'S BONE
 THE SPITFIRE GRILL
 ROOM IN ROME

Top Rated Most Popular

TV Shows

Mix-and-match from the categories below...

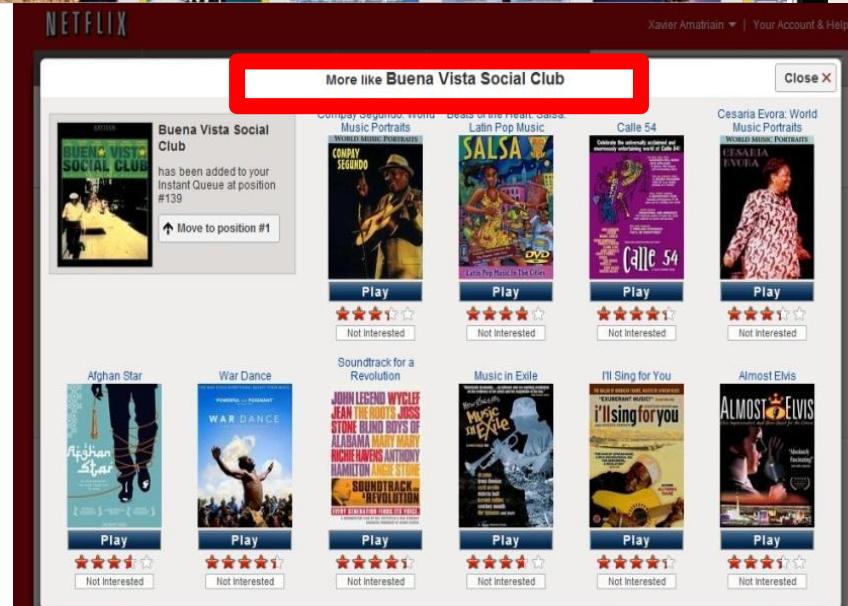
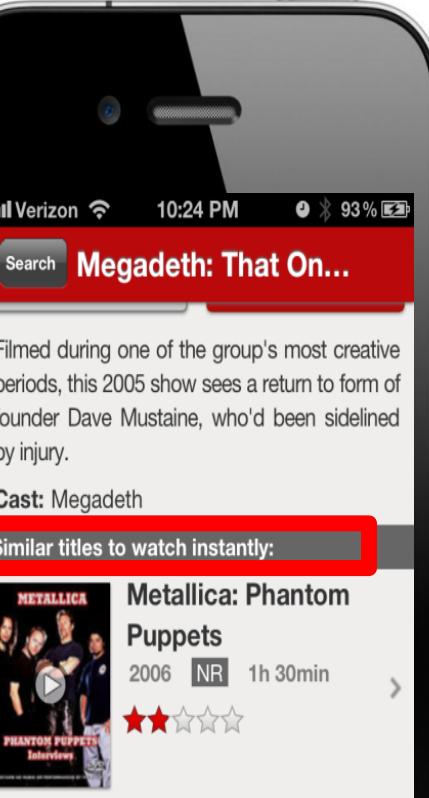
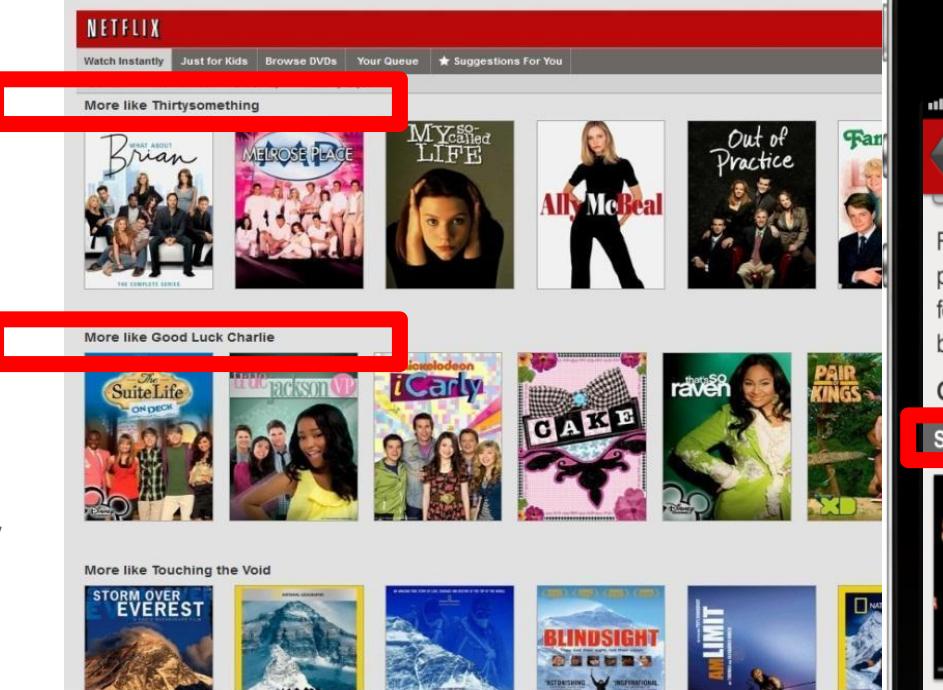
Family-friendly
 TV Comedies
 Cartoons
 Kids' TV Shows
 TV Docs

NE

Similar

- Displayed in many different contexts
 - In response to user actions/context (search, queue add...)
 - More like... rows



Page Composition

NETFLIX

Watch Instantly · Just for Kids · Taste Profile · DVDs · DVD Queue

Recently Watched · My List · See All

Popular on Netflix

Critically-acclaimed Movies

New Releases

Kids' TV

Because you watched Revolution

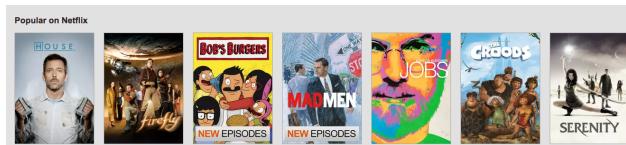
Because you watched Stargate

Top Picks for Justin

Children & Family Movies

Page Composition

10,000s
of
possible
rows



...



Variable number of
possible videos per
row (up to
thousands)

1 personalized
page



10-40
rows

per
device



Page Composition

Accurate vs. Diverse

Discovery vs. Continuation

Depth vs. Coverage

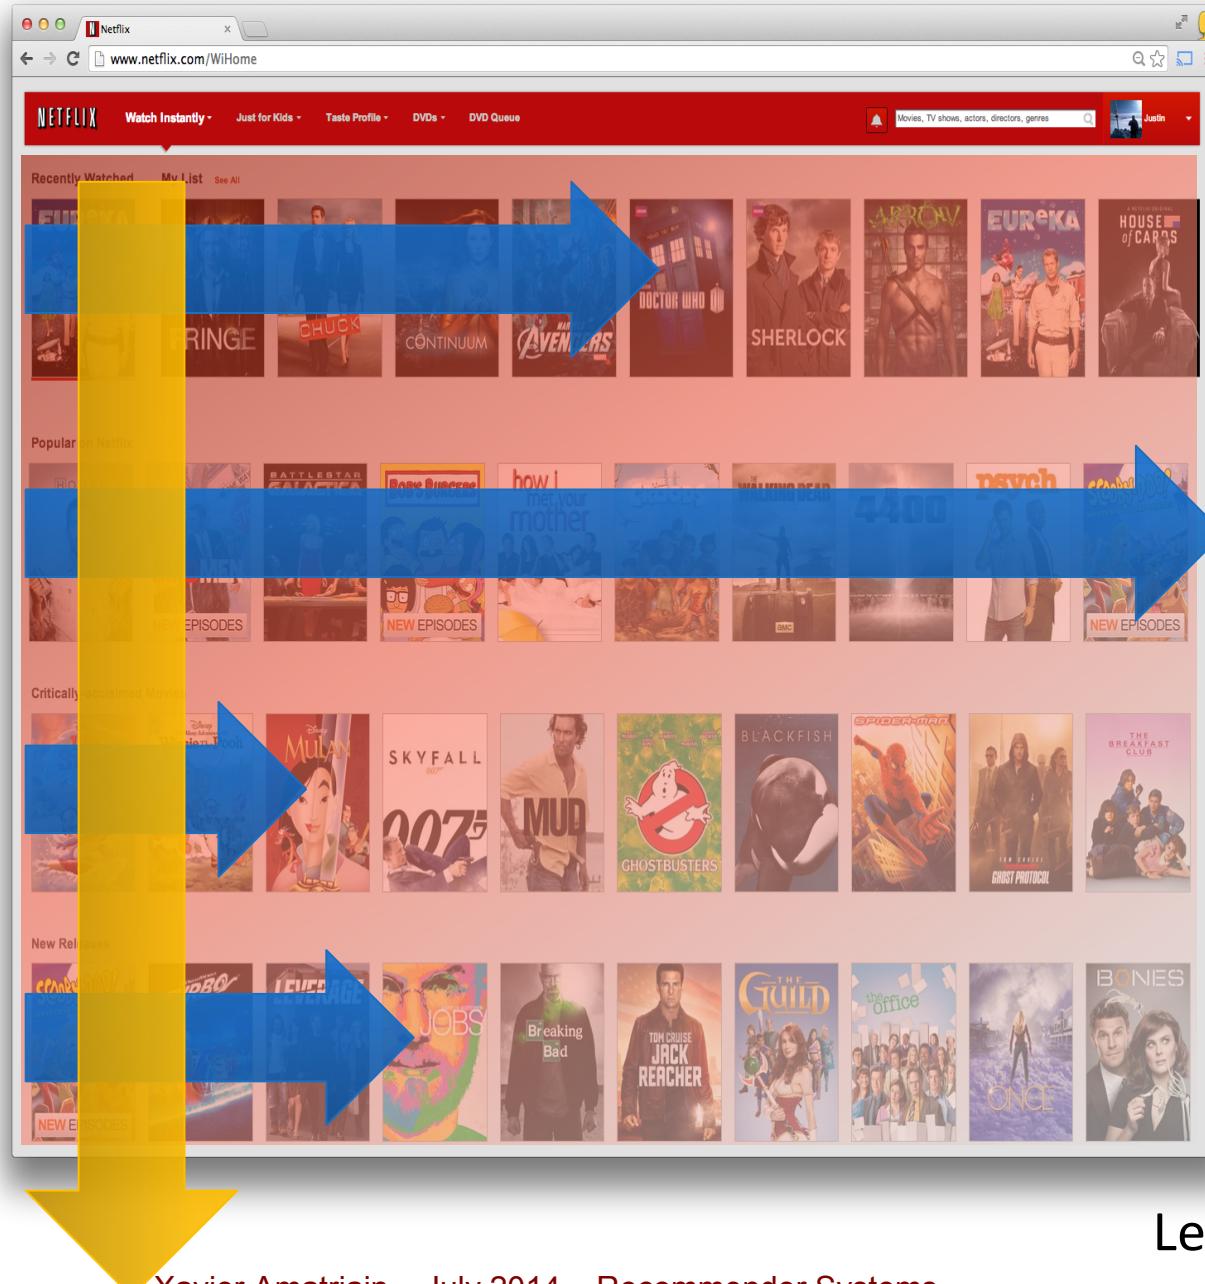
Freshness vs. Stability

Recommendations vs. Tasks



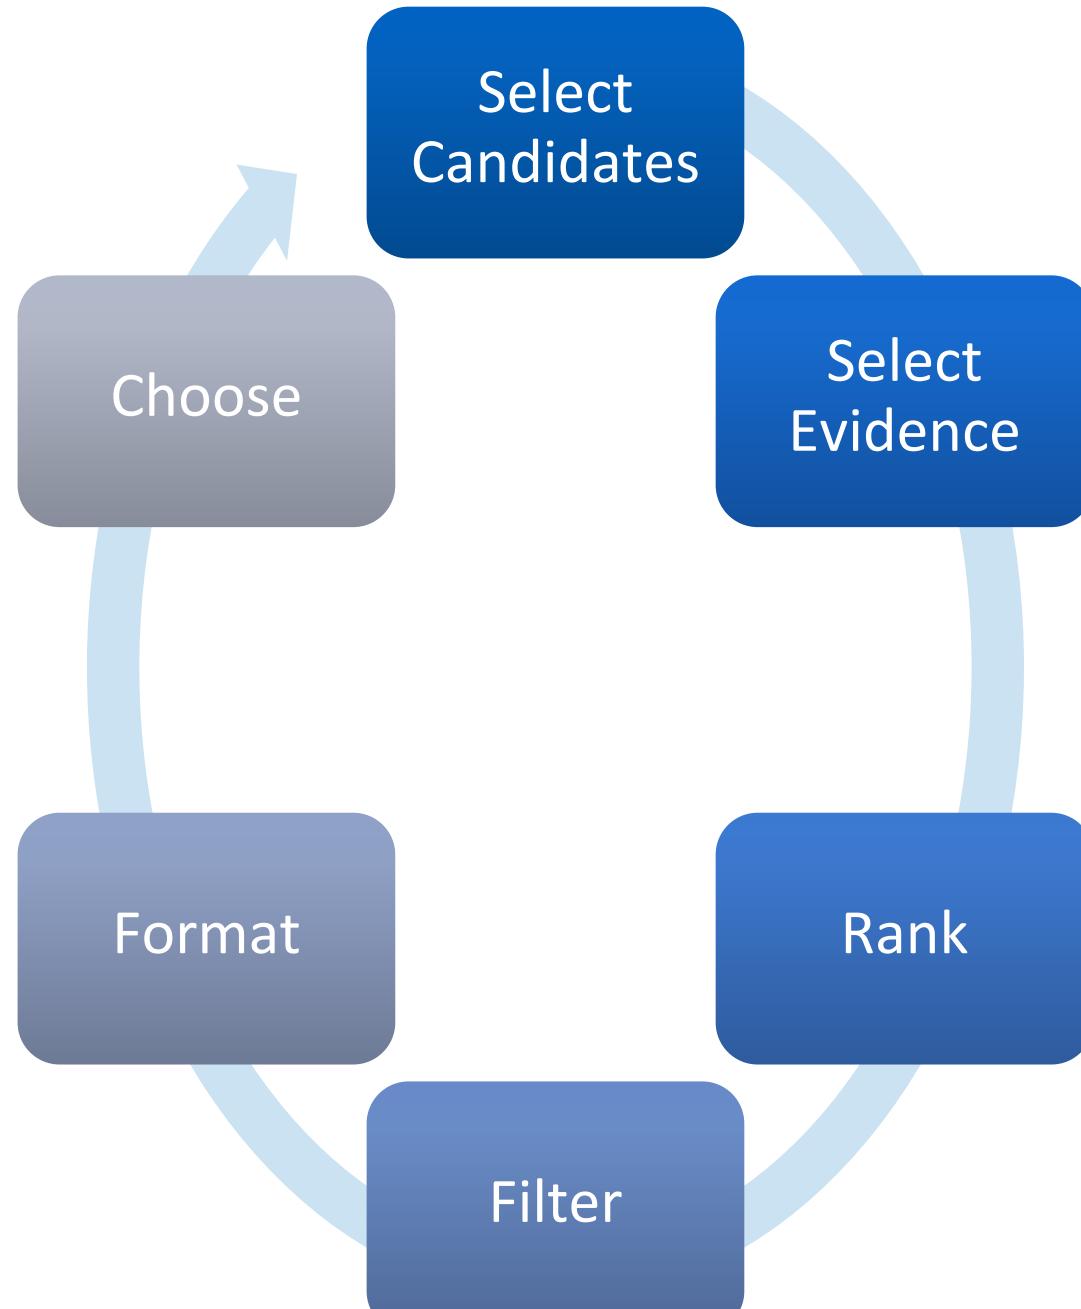
Page Composition

More likely to see



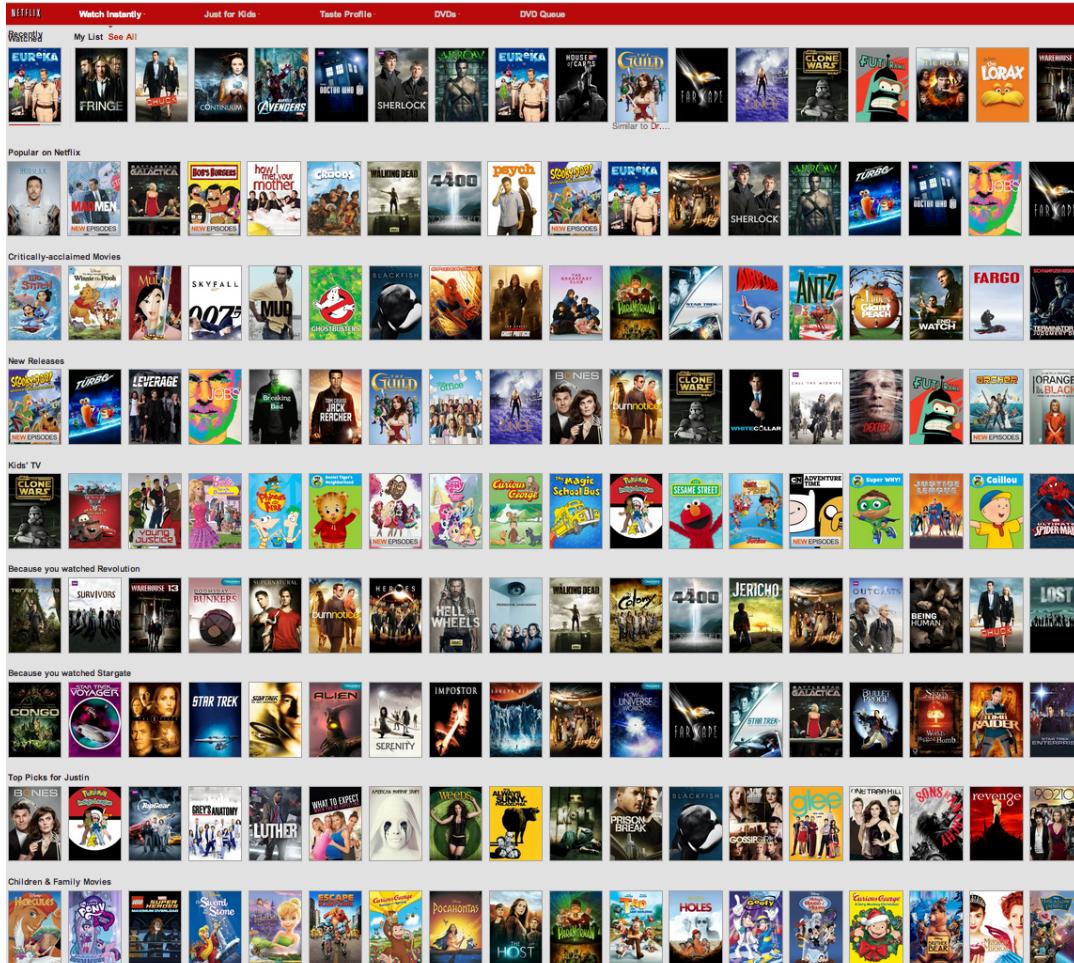
Less likely

Page Composition



Row Features

- Quality of items
- Features of items
- Quality of evidence
- User-row interactions
- Item/row metadata
- Recency
- Item-row affinity
- Row length
- Position on page
- Context
- Title
- Diversity
- Freshness
- ...



Search Recommendations

MOVIES & TV SHOWS

Marriage Italian Style

Smalltown, Italy

Perlasca

PEOPLE

Italia Almirante-Moreira

Italo Renda

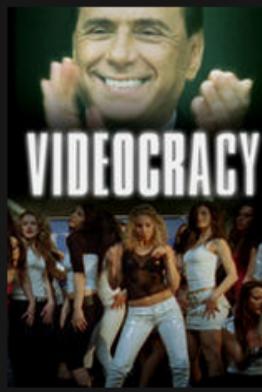
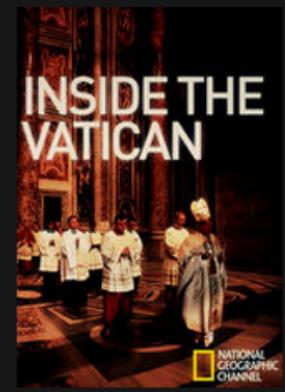
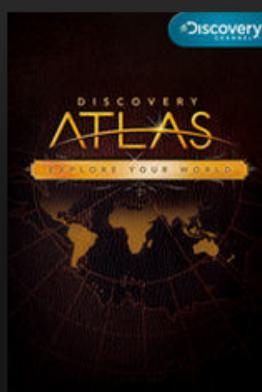
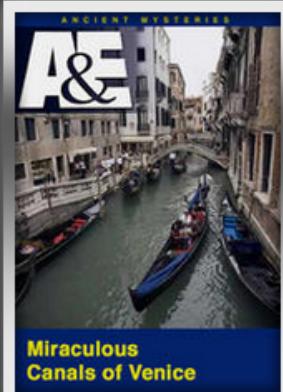
EXPLORE TITLES RELATED TO ITALY

The Italian Job

The Italian

italy

Related to italy



Ancient Mysteries: Canals of Venice

2005 TV-G 46m

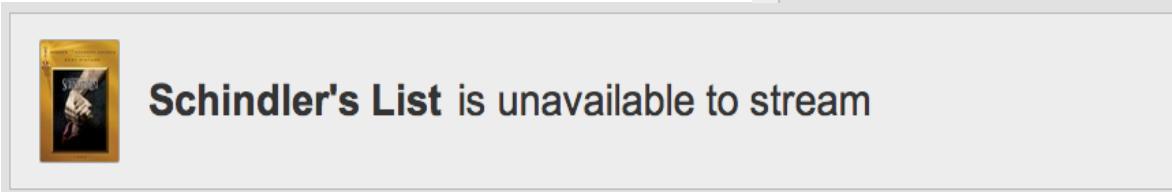


Known for its distinctive man-made canals and unparalleled aura of romance, the Italian city of Venice is like no other place on Earth.

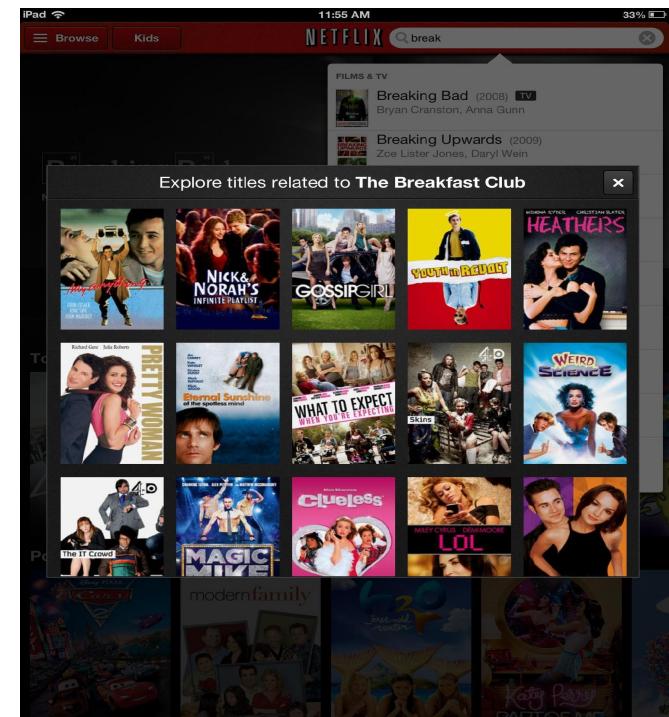
TV Shows, Documentaries

Unavailable Title Recommendations

The screenshot shows a search results page for "Schindler's List". The search bar contains the query. Below it, there are two main sections: "PEOPLE" and "UNAVAILABLE". Under "UNAVAILABLE", several titles are listed: "Dagmar Schindel-Hildebrand", "Schindler's List", "Schindler's Houses", "Oskar Schindler: The True Story", and "Varian Fry: Artists' Schindler".



After searching for this title, many members stream:



Search Recommendations

Combination of Different sources

- Transition on Play after Query
- Similarity between user's (query, play)
- Editorial
- ...

Metaphor: a System for Related Search Recommendations

Azarias Reda
University of Michigan
azarias@umich.edu

Yubin Park
University of Texas at Austin
yubin.park@utexas.edu

Mitul Tiwari
LinkedIn
mtiwari@linkedin.com

Christian Posse
LinkedIn
cposse@linkedin.com

Sam Shah
LinkedIn
samshah@linkedin.com



Postplay

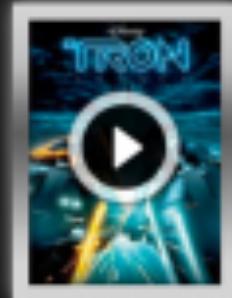


Rate The Fighter



Back to browse Search

Suggestions for you
to watch now...



A thumbnail image for the movie Tron: Legacy, showing Sam Flynn in his Tron suit riding a lightcycle through a futuristic digital environment. The poster includes the title "Tron: Legacy", the rating "PG 2h 05m", and a 5-star rating icon.

NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

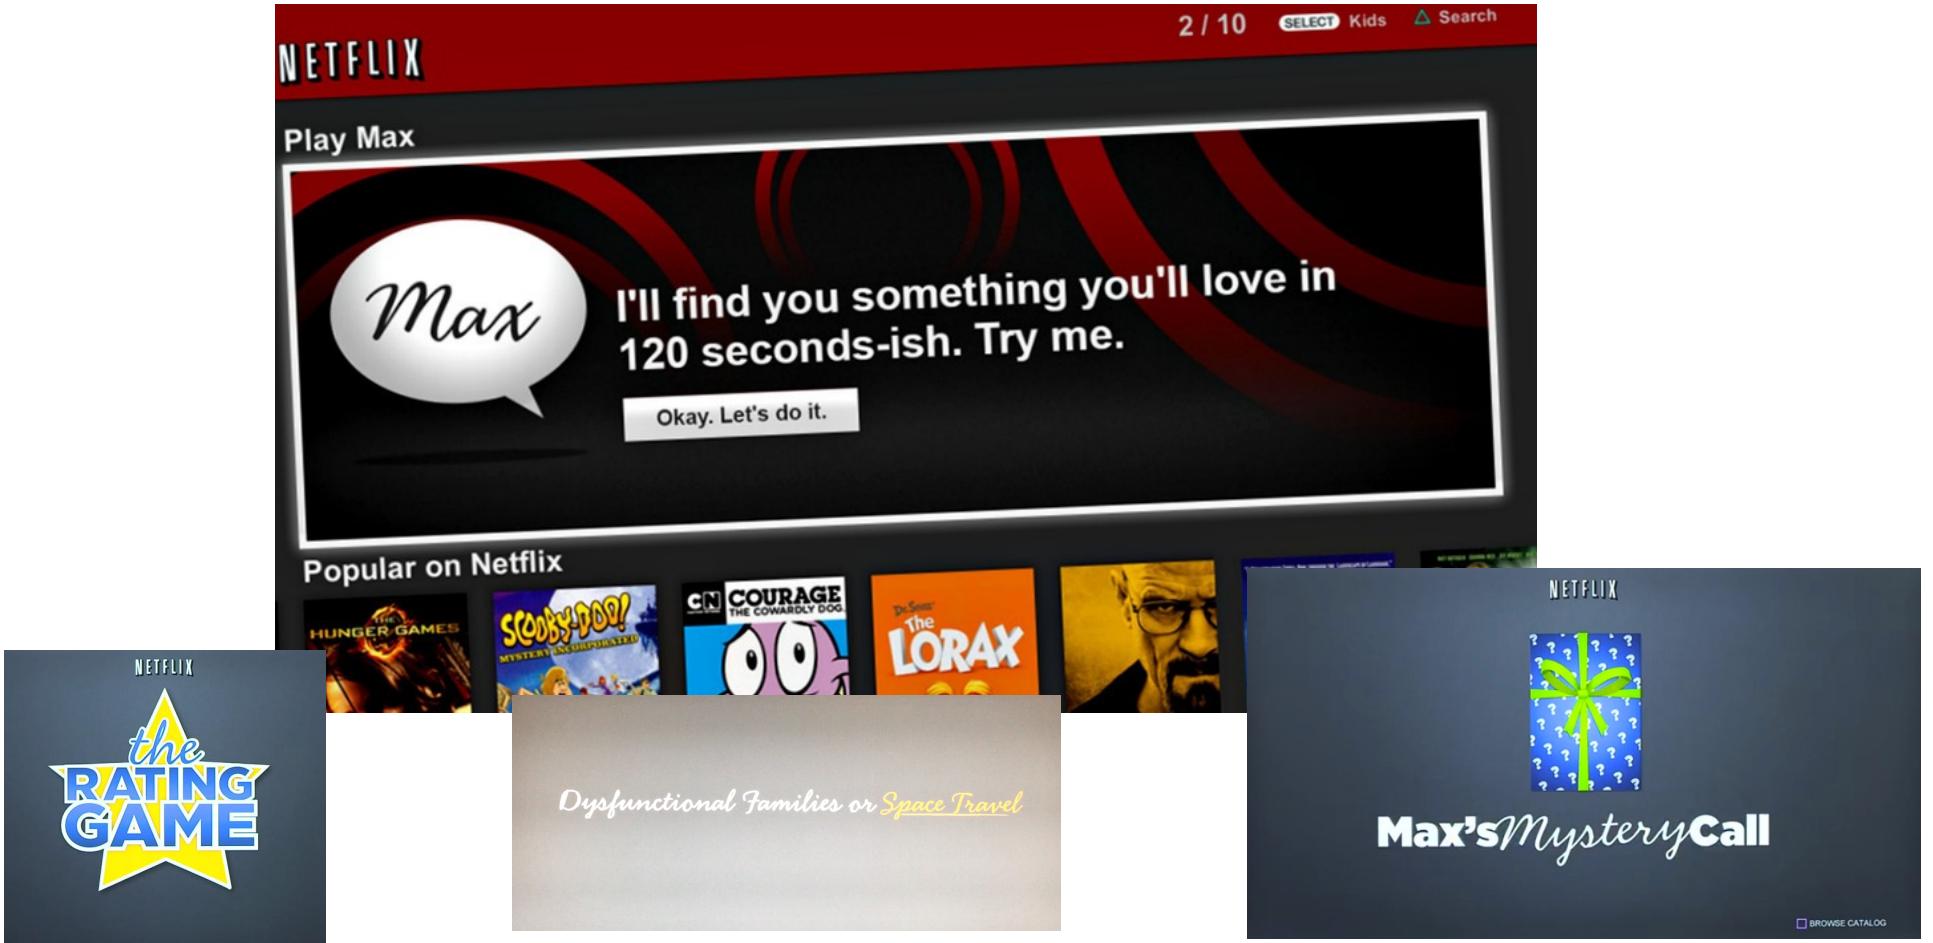
Billboard



NETFLIX

Xavier Amatriain – July 2014 – Recommender Systems

Gamification



NETFLIX

Diversity & Awareness

Personalization awareness



6. Conclusions

Conclusions

- For many applications such as Recommender Systems (but also Search, Advertising, and even Networks) understanding data and users is vital
- Algorithms can only be as good as the data they use as input
 - But the inverse is also true: you need a good algorithm to leverage your data
- Importance of User/Data Mining is going to be a growing trend in many areas in the coming years

Conclusions

- Recommender Systems (RS) is an important application of User Mining
- RS have the potential to become as important as Search is now
- However, RS are more than User Mining
 - HCI
 - Economical models
 - ...

Conclusions

- RS are fairly new but already grounded on well-proven technology
 - Collaborative Filtering
 - Machine Learning
 - Content Analysis
 - Social Network Analysis
 - ...
- However, there are still many open questions and a lot of interesting research to do!